



---

# Module 13 – Main Memory

---

**(ECE M116C- CS M151B) Computer Architecture Systems**

**Nader Sehatbakhsh**


**Department of Electrical and Computer Engineering**

**University of California, Los Angeles**

# Last time...

- Use memory hierarchy to achieve both *fast* and *large* memory system.
  - Cache performance: hit time, miss rate, and miss penalty.
  - Replacement policy
  - How to improve each metric

# Main Memory




# Technology

- We see combination of **DRAM** and **DISK** as the main memory (usually call the DRAM part Main Memory and the HardDrive part DISK).
- Access to the Main Memory (combination of disk and DRAM) *is always a hit*.
  - As we will describe later, data could be in disk and/or DRAM, and we need to handle that.
  - Huge difference between latency of DRAM to DISK (and of course to cache).

# Where to store different programs?


→ *How to run different programs in a computer?*

- Fundamental Requirement: *Isolation*



# Where to store different programs?

- Each program typically has *different regions*:
  - code
  - data
  - heap
  - stack
  - ...
- Portion of the memory code be **shared** among multiple programs.
- Portion of the memory can be **reserved** for the OS and/or other privileged activities.



# Memory Management

- Early (and simple) machines only run one program with unrestricted access to ALL memory.
- Modern systems run several programs, and memory should be shared between multiple programs.


- Memory management is controlled by OS or system software, or hardware.

# How to share the memory between multiple programs?



# How to share the memory between multiple programs?

★ *Option 1: partition the memory (statically)!*




# How to share the memory between multiple programs?

## ★ *Option 1: partition the memory (statically)!*

### ○ Considerations:

- **Protection:** Independent program should not be able to affect each other inadvertently.
- **Location-Independence:** Program might be moved anywhere in the memory.




# How to share the memory between multiple programs?

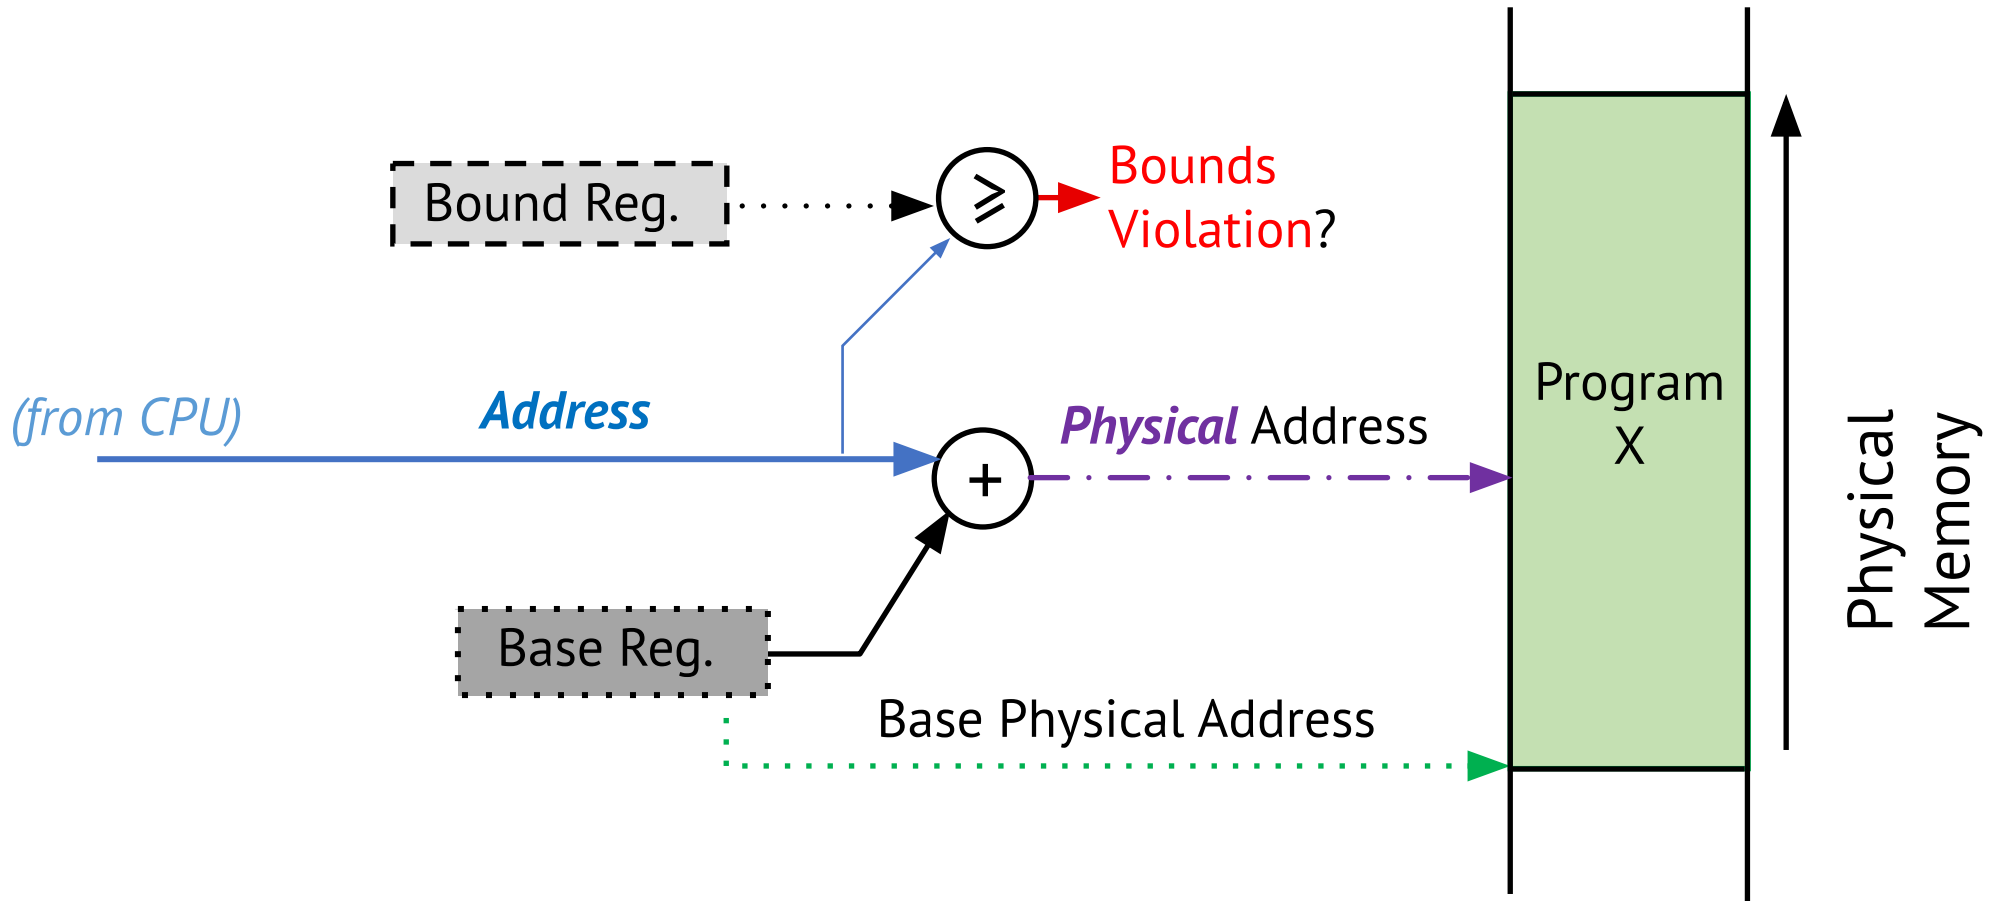
## ★ *Option 1: partition the memory (statically)!*

### ○ Considerations:

- **Protection:** Independent program should not be able to affect each other inadvertently.
  - *check bounds!*
- **Location-Independence:** Program might be moved anywhere in the memory.
  - *use a base pointer!*



# Static Partitioning



# Address Translation

- To access the memory, addresses in a program (typically called logical or effective address) should be translated into a physical address.


Example:

$$PA = Base\_Reg + Logical\_Adr$$

- When switching programs, OS updates base and bound registers.

# How to share code/libraries?


- Multiple programs can access (read-only) shared memory spaces.
  - Standard libraries (e.g., printf)
  - Drivers
  - etc.



# How to share code/libraries?

- Multiple programs can access (read-only) shared memory spaces.
  - Standard libraries (e.g., printf)
  - Drivers
  - etc.


❖ *What about the OS?*



# How to share code/libraries?

- Multiple programs can access (read-only) shared memory spaces.
  - Standard libraries (e.g., printf)
  - Drivers
  - etc.


❖ *What about the OS?*  
*- It's a program as well!*





# How to share the memory between multiple programs?


★ *Option 1: partition the memory (statically)!*



# How to share the memory between multiple programs?

★ *Option 1: partition the memory (statically)!*


*- What happens if we add a new program?*



# How to share the memory between multiple programs?

★ *Option 1: partition the memory (statically)!*

- *What happens if we add a new program?*
- *What if one program need more (less) space?*



# How to share the memory between multiple programs?

★ *Option 2:*

# How to share the memory between multiple programs?

★ *Option 2: partition the memory (dynamically)!*

# How to share the memory between multiple programs?

- ★ *Option 2: partition the memory (dynamically)!*
  - Divide the memory into fixed-size blocks (called page).

# How to share the memory between multiple programs?

- ★ *Option 2: partition the memory (dynamically)!*
  - Divide the memory into fixed-size blocks (called page).

0
1
2
3

Program A

0
1
2
3

Program B

0
1
2
3

Program C

# How to share the memory between multiple programs?

- ★ *Option 2: partition the memory (dynamically)!*
  - Divide the memory into fixed-size blocks (called page).

0
1
2
3

Program A

-----> Each page is typically 4KB. For example, we can use the lower 12 bits of each address to compute its page number.



# How to share the memory between multiple programs?

- ★ *Option 2: partition the memory (dynamically)!*
  - Divide the memory into fixed-size blocks (called page).

0
1
2
3

Program A

0
1
2
3

Program B

0
1
2
3

Program C

1
0
1
3
3
2
0
0
2
2
3
1
Operating System Pages

# How to share the memory between multiple programs?

## ★ *Option 2: partition the memory (dynamically)!*

- Divide the memory into fixed-size blocks (called page).
- Assign pages to each program as needed.
- Pages can be *scattered* in the memory.
- Typically 4KB per page.

0
1
2
3

Program A

0
1
2
3

Program B

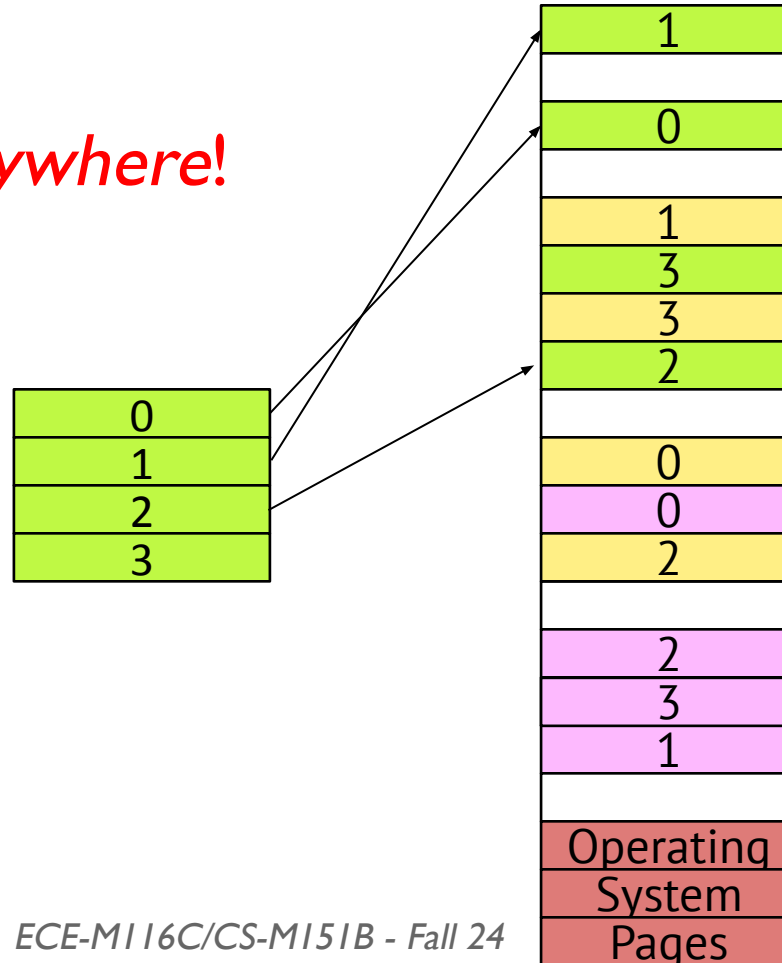
0
1
2
3

Program C

1
0
1
3
3
2
0
0
2
2
3
1
Operating System Pages

# How to track which page is where?

- Each page could be *anywhere*!




# How to track which page is where?

- Each page could be *anywhere*!

-- Use a table to keep track of mapping!

0
1
2
3


0	p0
1	p1
2	p2
3	p3



# How to track which page is where?

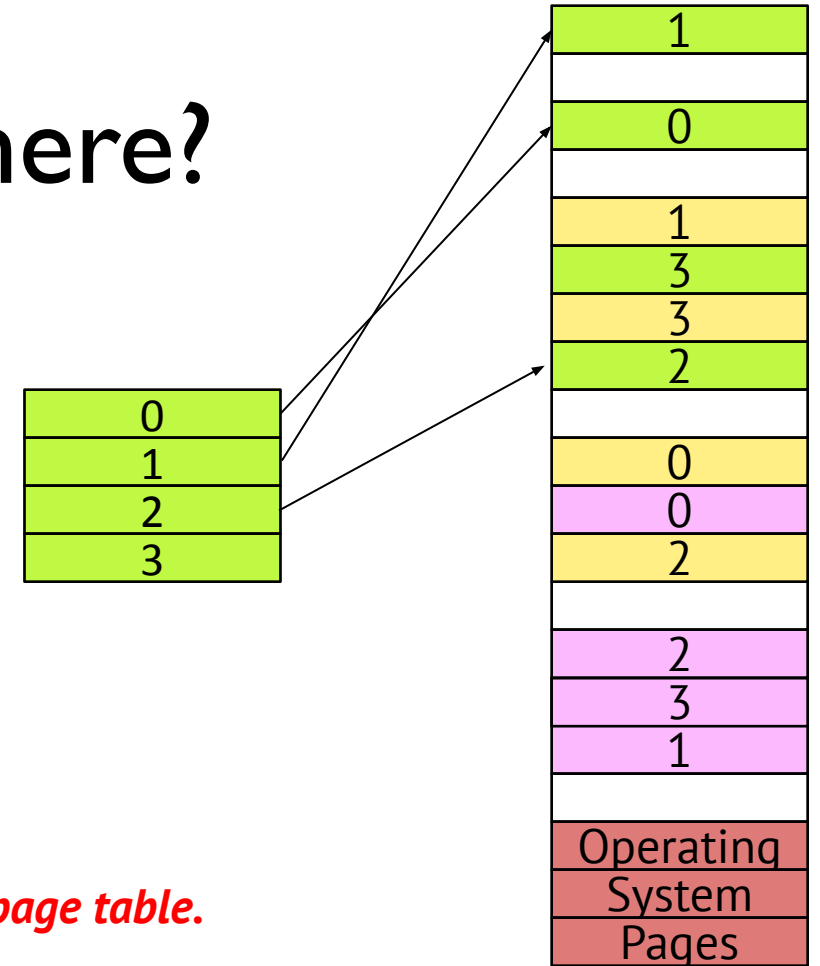
- Each page could be *anywhere*!

-- Use a table to keep track of mapping!



0	p0
1	p1
2	p2
3	p3

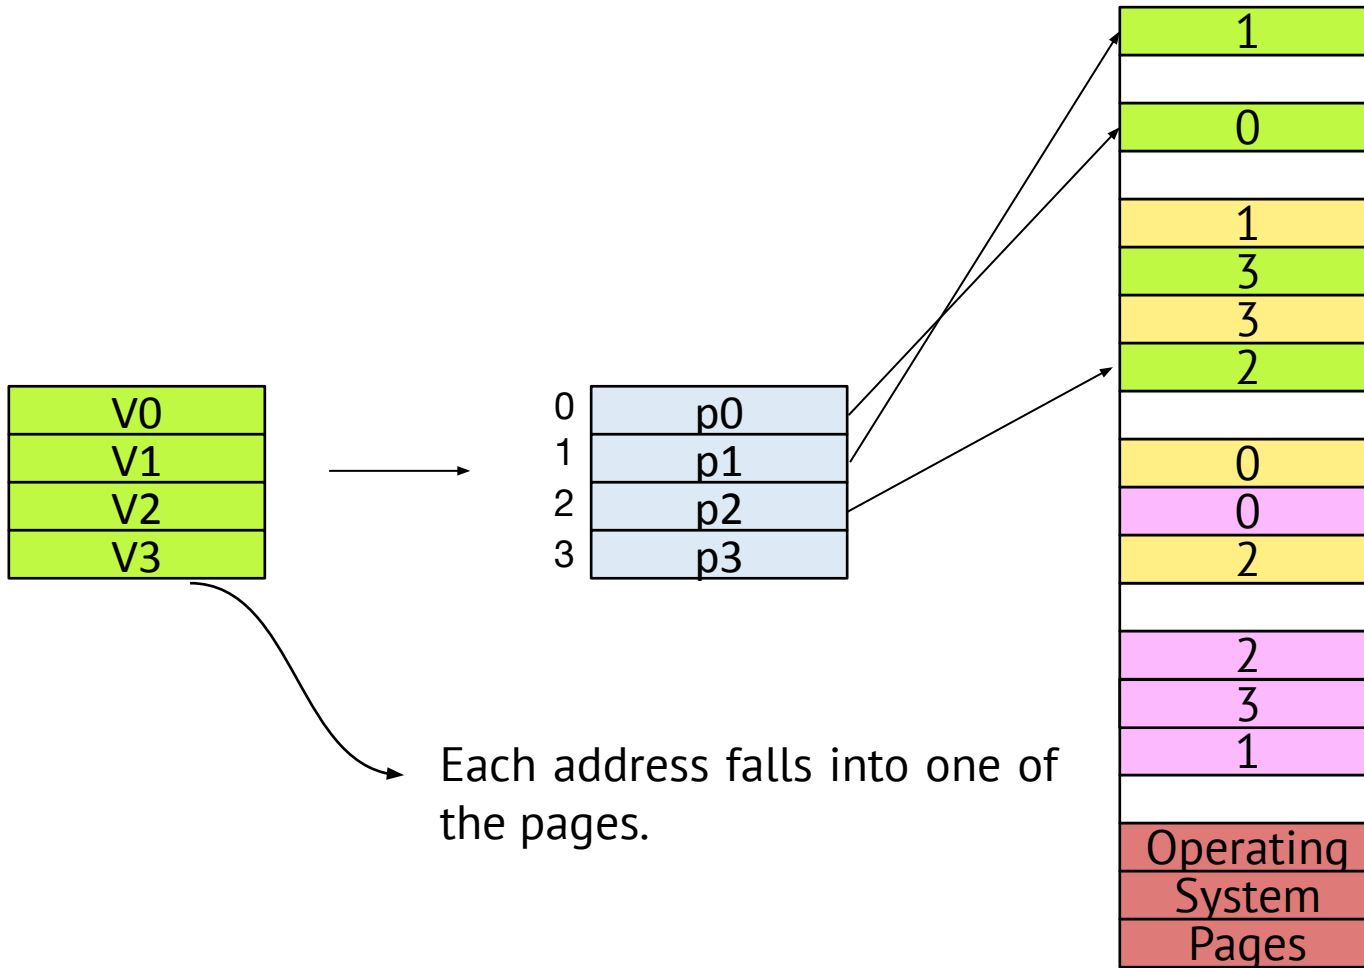
→ We call this a *page table*.



# Virtual Address

- Since we are using page tables for translations, we no longer need to use real physical memories addresses in each program.
  - Each program can start from (virtual) address 0x00.
  - Each program sees a *large, private, and uniform* memory.  
*(it's just an illusion though!)*

# Virtual to Physical Address Translation



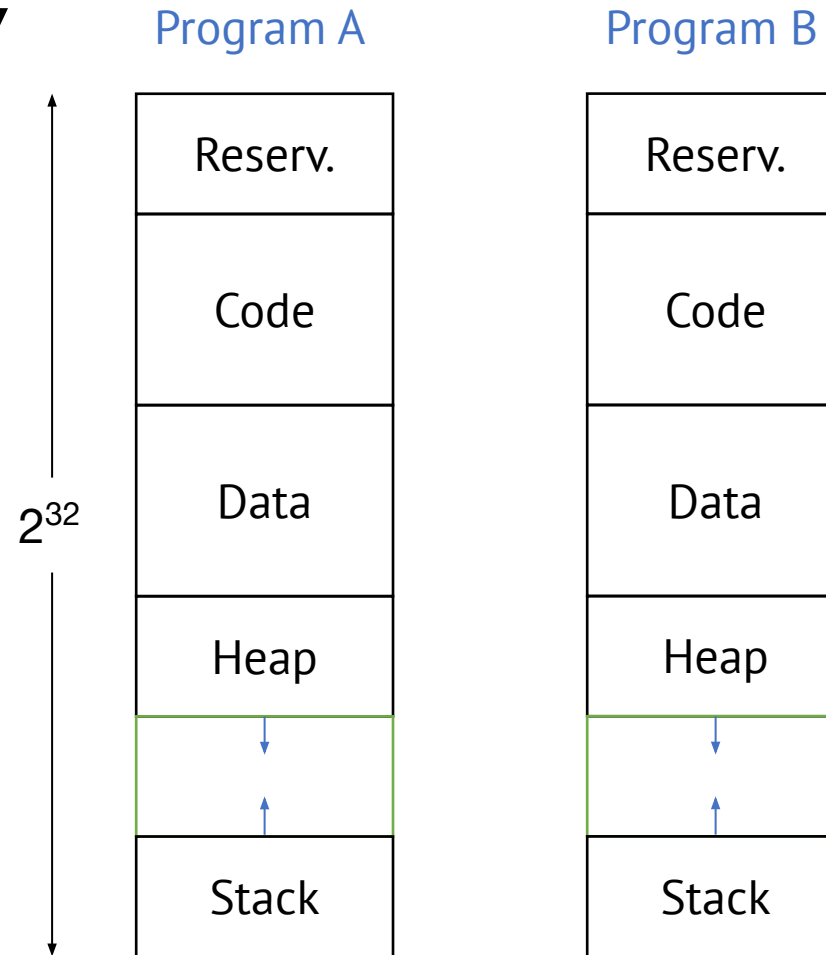
# Address Lifecycle

- Program → VA → PA → Memory



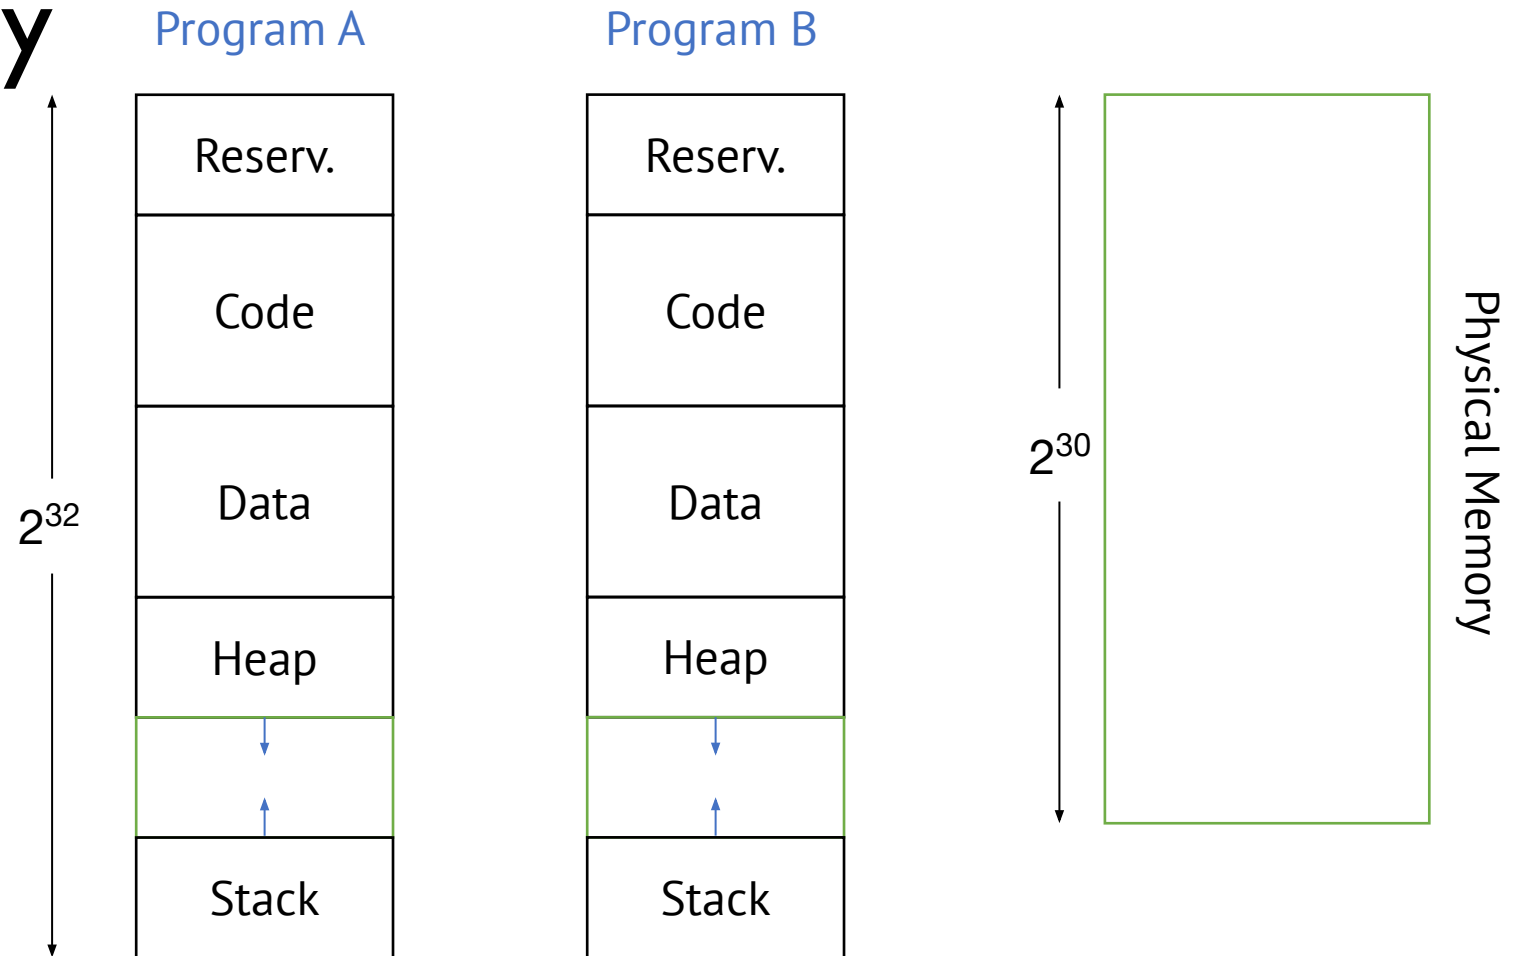
# Virtual Memory

*Private, large, and unified!*

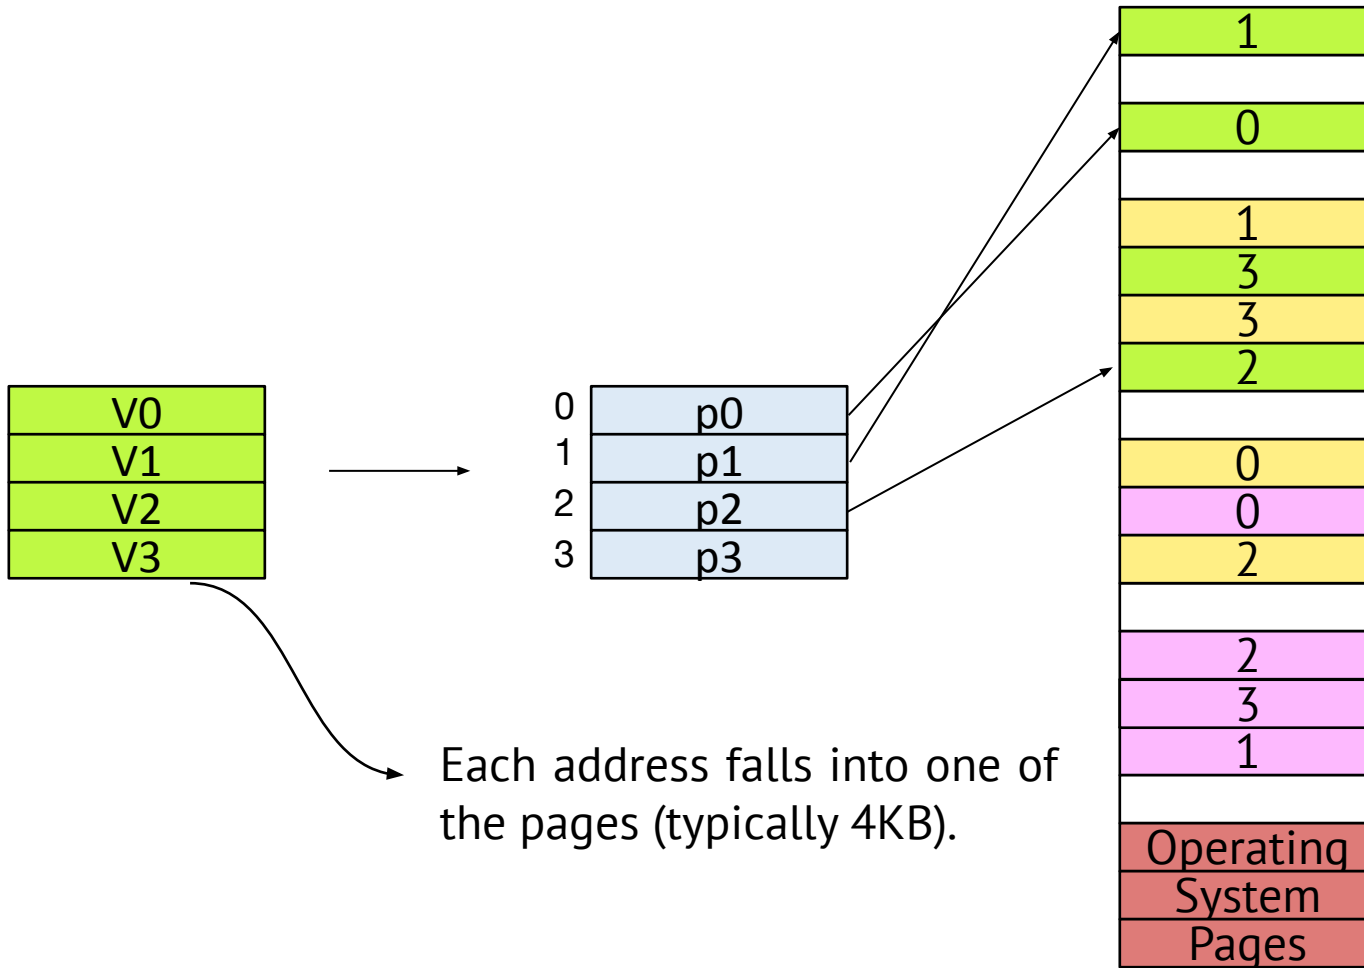


# Physical Memory

*What we actually have!*



# Virtual to Physical Address Translation

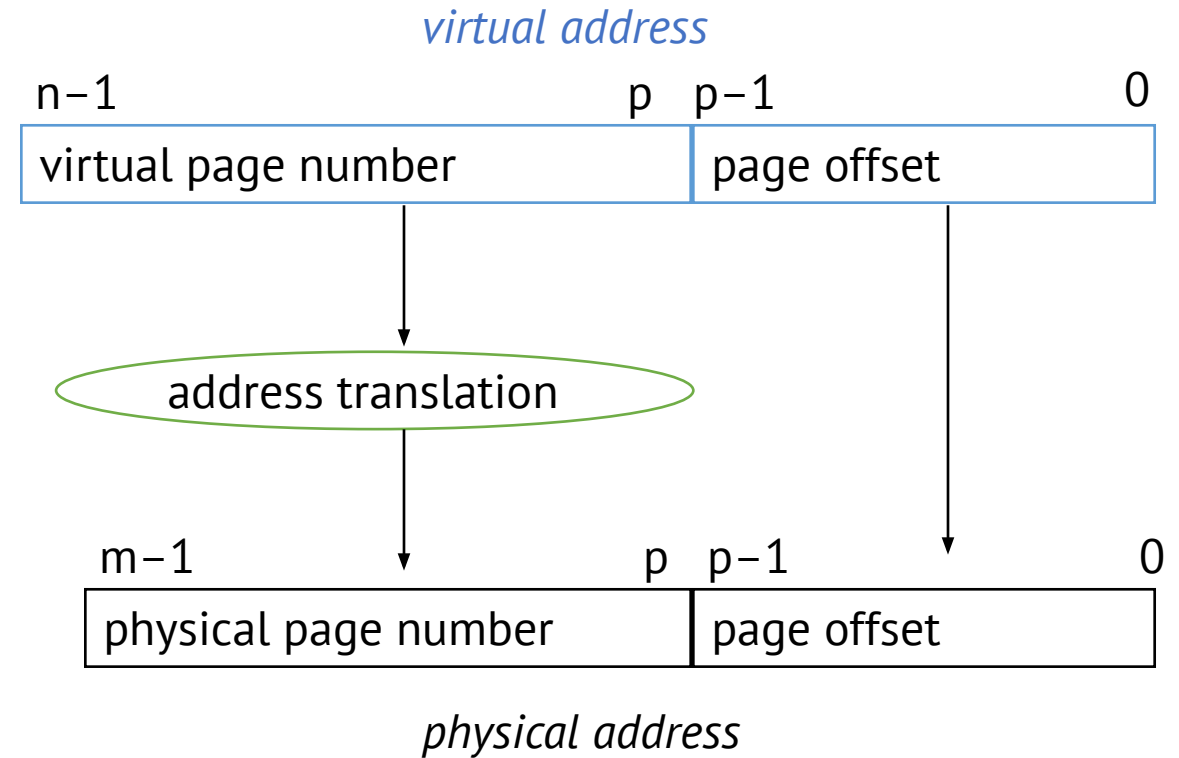


# Virtual to Physical Address Translation

$2^n$ : virtual address size

$2^m$ : physical address size

$2^p$ : page size



# Virtual and Physical Memory

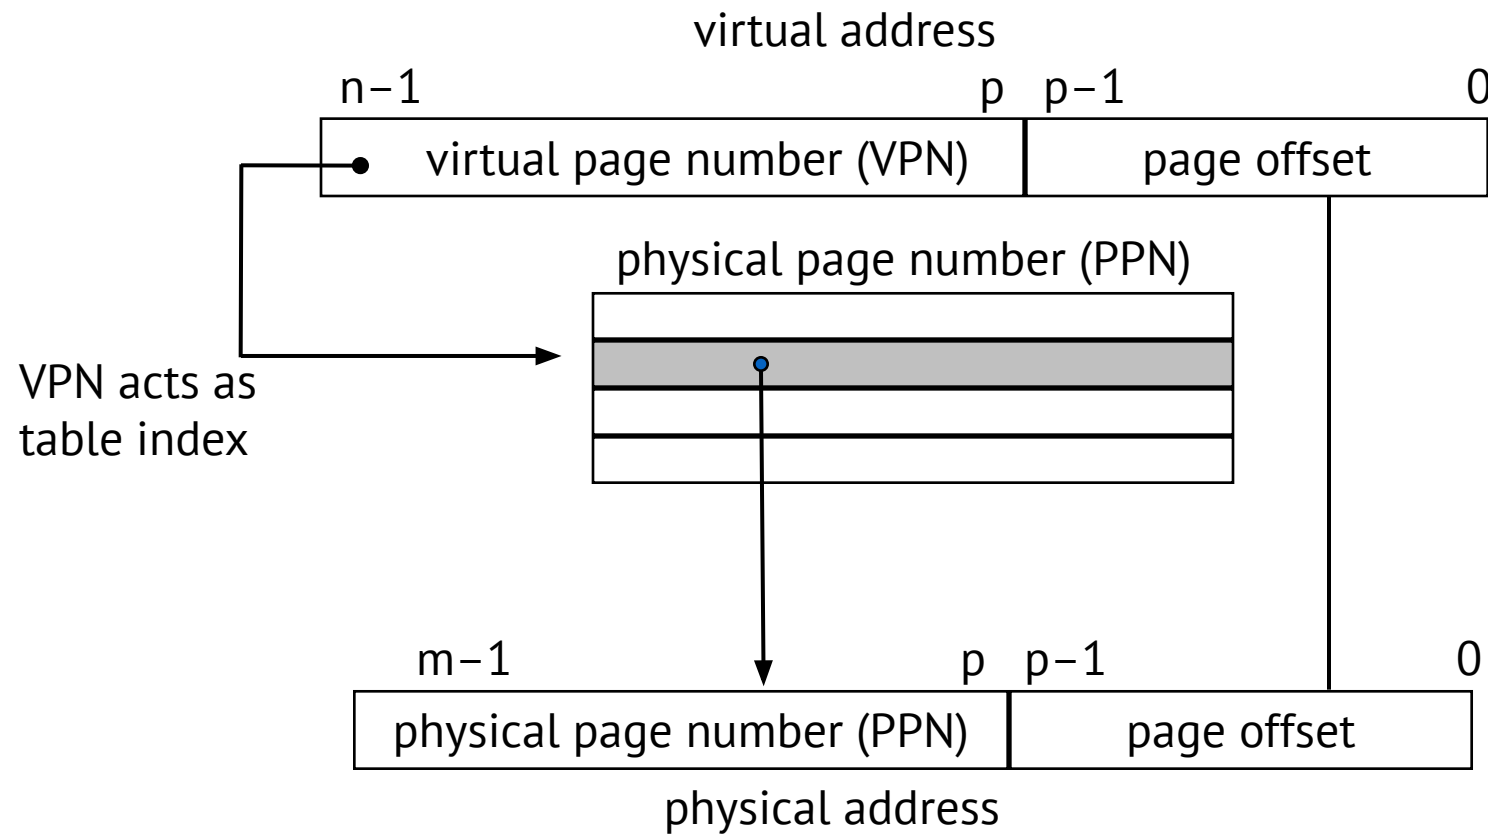
- System:

- Virtual memory size: 4 GB =  $2^{32}$  bytes
- Physical memory size: 256 MB =  $2^{28}$  bytes
- Page size: 4 KB =  $2^{12}$  bytes

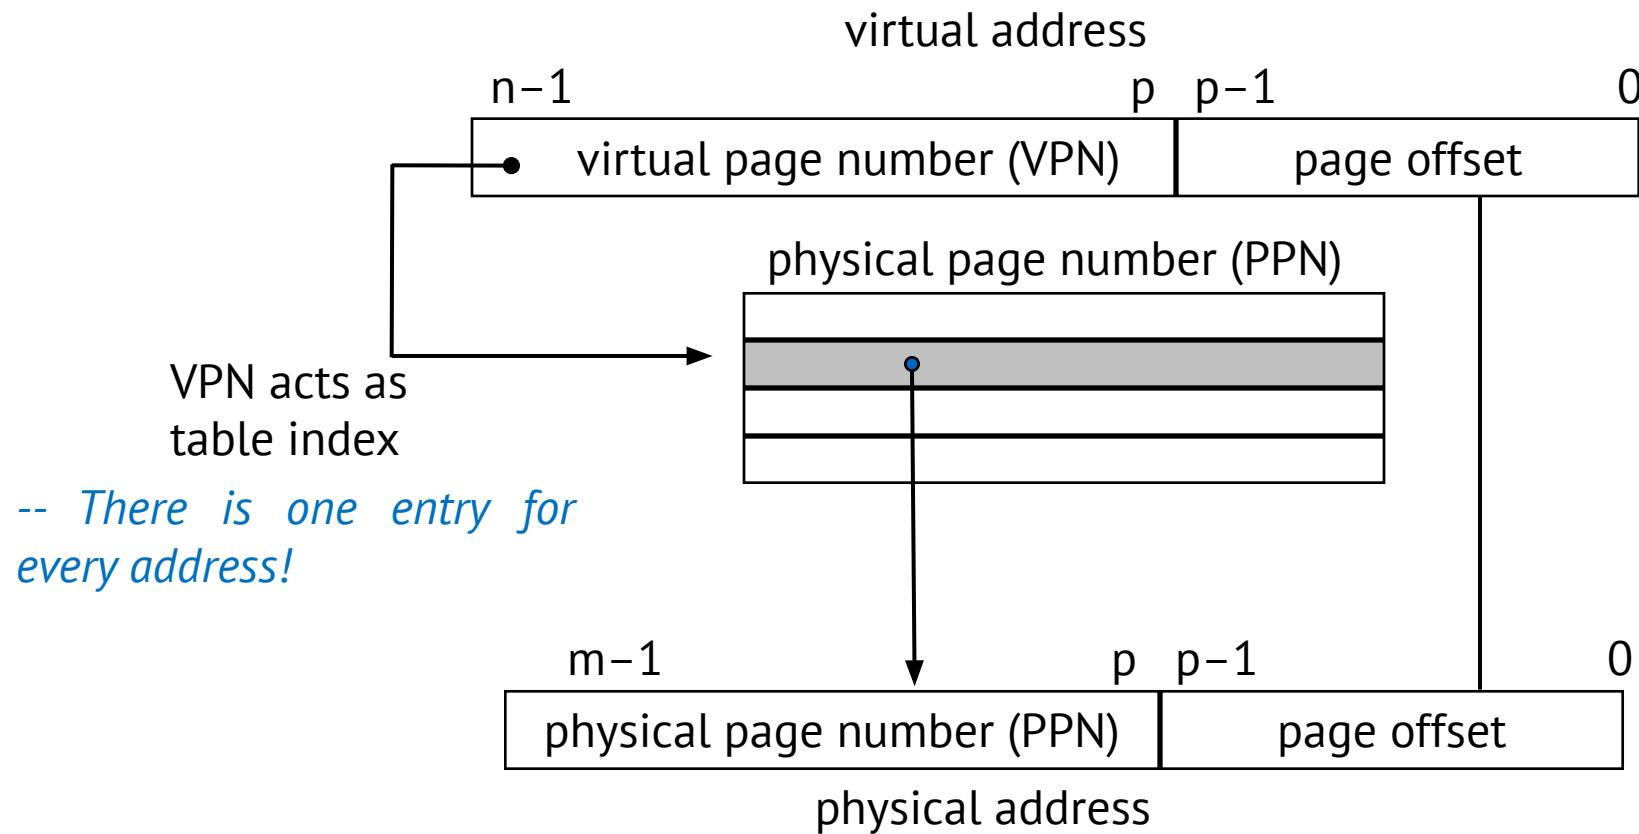
- Organization:

- Virtual address: 32 bits
- Physical address: 28 bits
- Page offset: 12 bits
- # Virtual pages =  $2^{32}/2^{12} = 2^{20}$  (VPN = 20 bits)
- # Physical pages =  $2^{28}/2^{12} = 2^{16}$  (PPN = 16 bits)

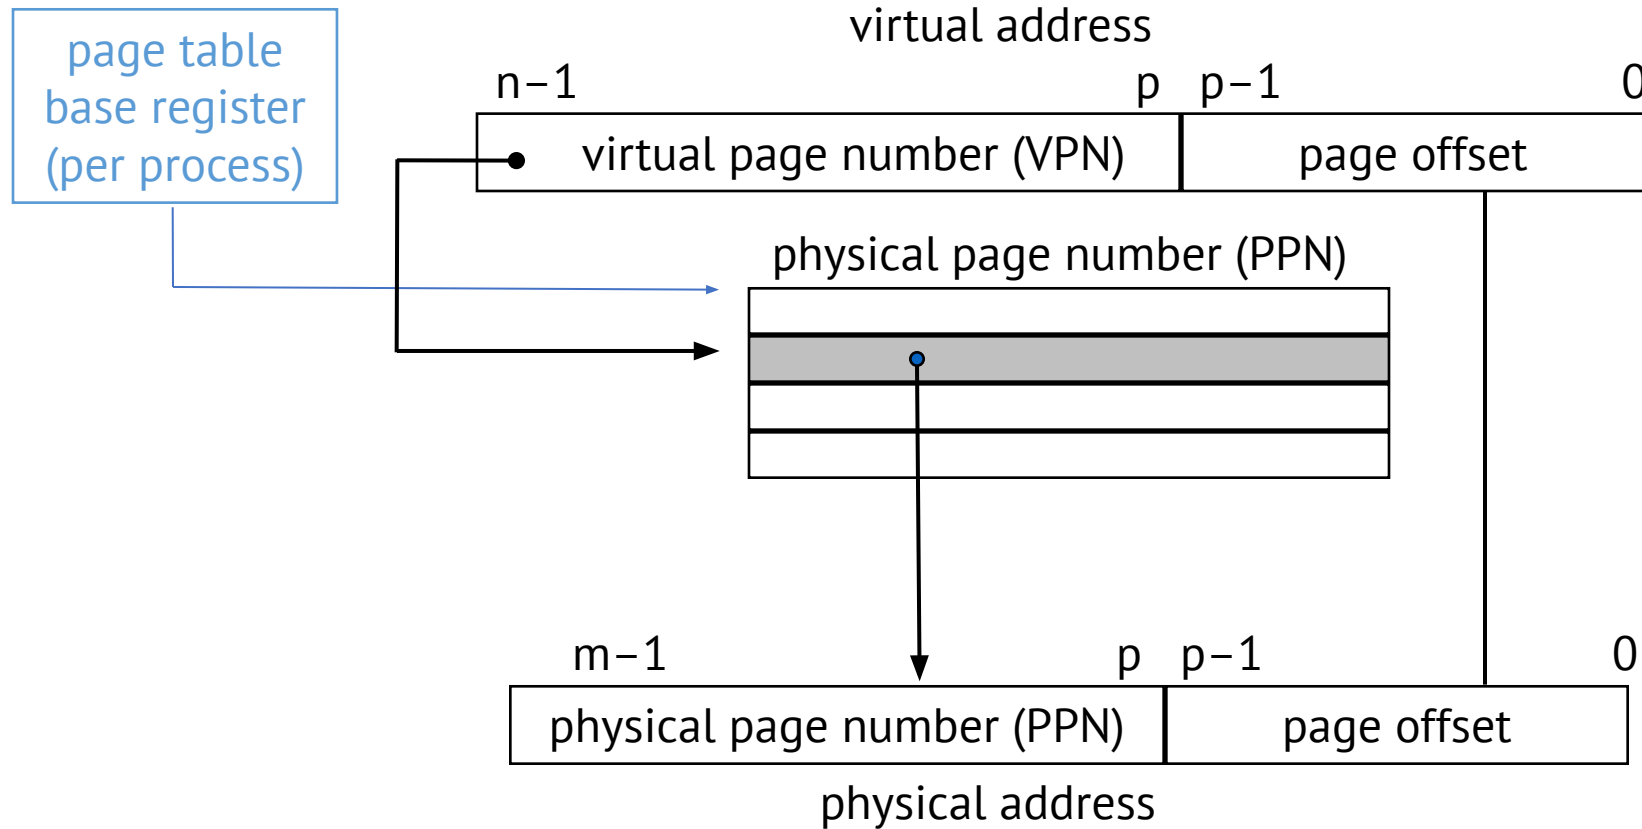
# Page Table



# Page Table



# Page Table





# What else to store in each page table entry (PTE)?

- Protections and flags
  - Whether we can write to this page or not (read-only)?
  - Who is the owner? Is it shared?
  - Can we execute it?

Flag	physical page number (PPN)

# Where to store page tables?

- Inside the memory itself.



# Where to store page tables?

- Inside the memory itself.
- ★ Each memory access becomes *two accesses*, one for accessing the page table (i.e.,  $\text{adr} = \text{VPN} + \text{base}$ ). The other for accessing the translated address (i.e.,  $\text{adr} = \text{PPN} + \text{page offset}$ ).

# How large is a page table?

- #Virtual pages =  $2^{32}/2^{12} = 2^{20}$  (VPN = 20 bits)
- Physical address: 28
- Each entry in page table (PTE) = 28 - offset + flags  $\approx$  32 bits (4 bytes)

→  $2^{20} \times 32 = 4 \text{ MB}$  for each application!!

(What about 64-bit address?)

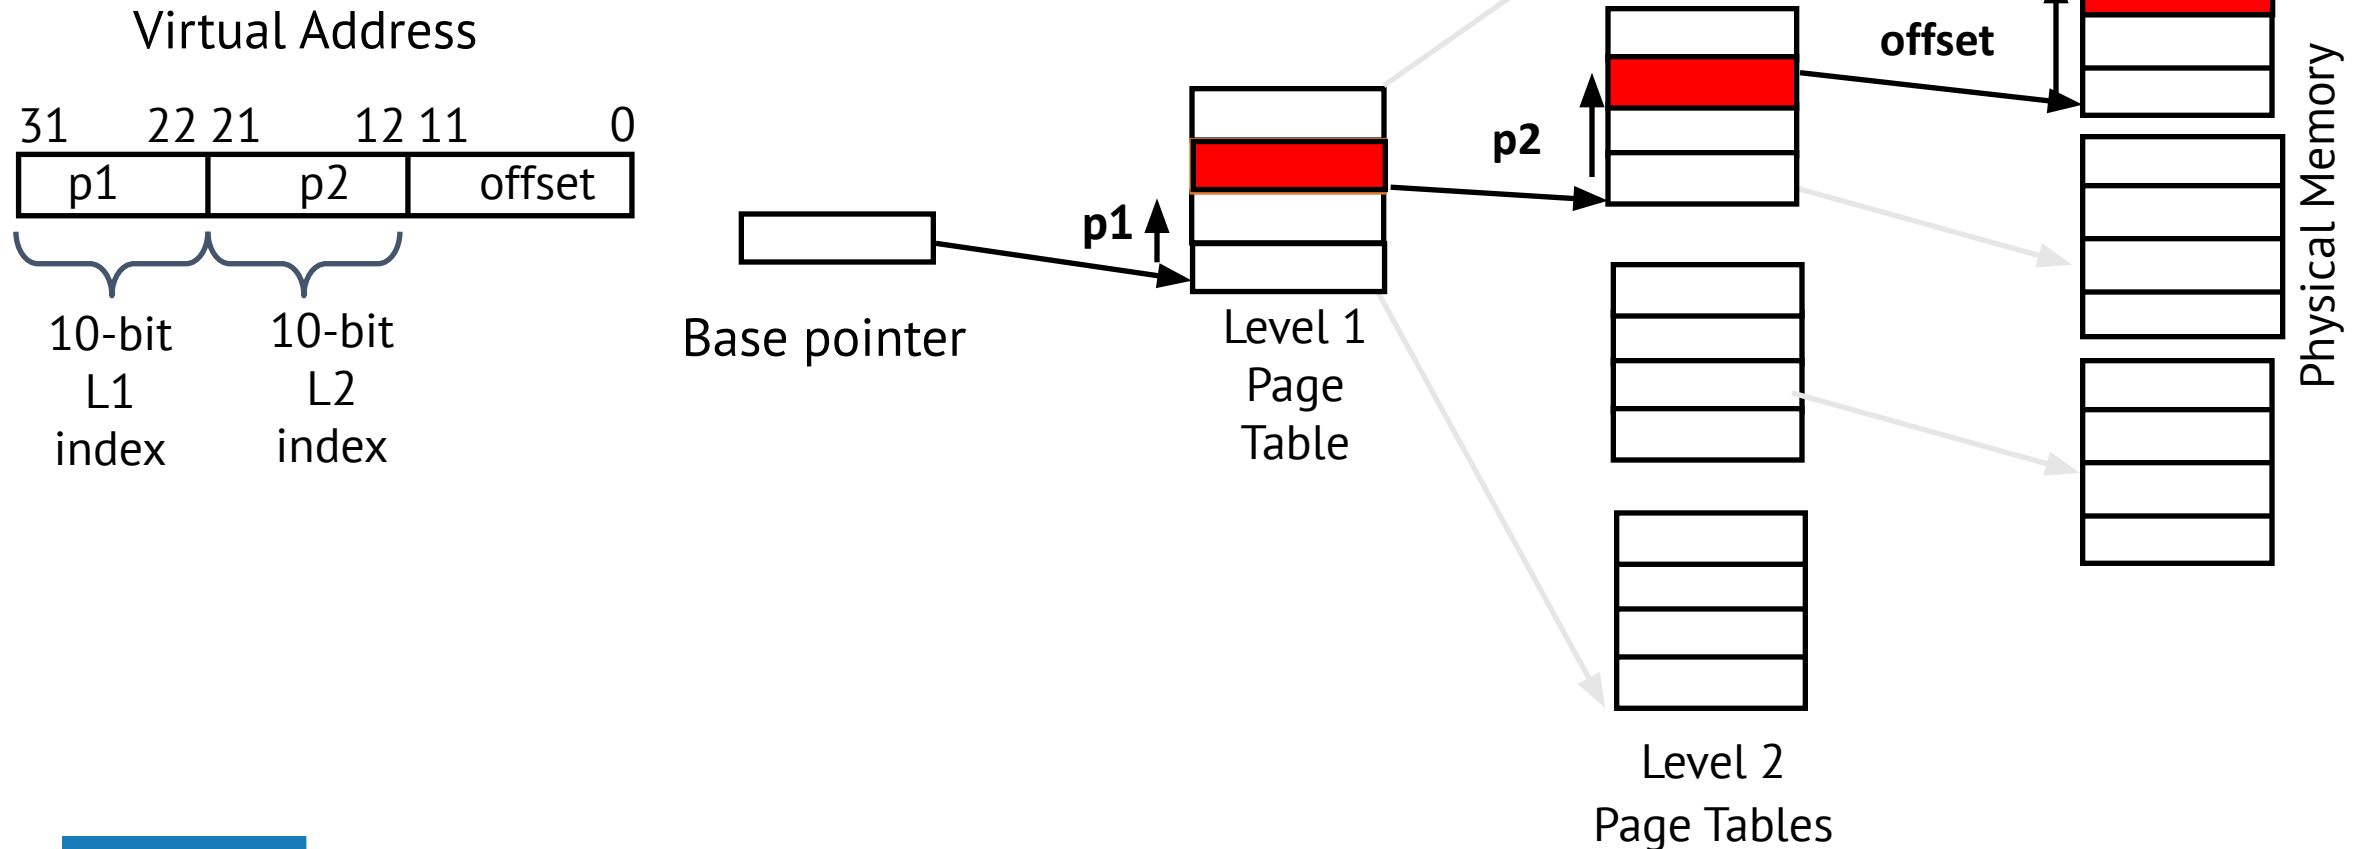
# How to compress page tables?

# Hierarchical Page table

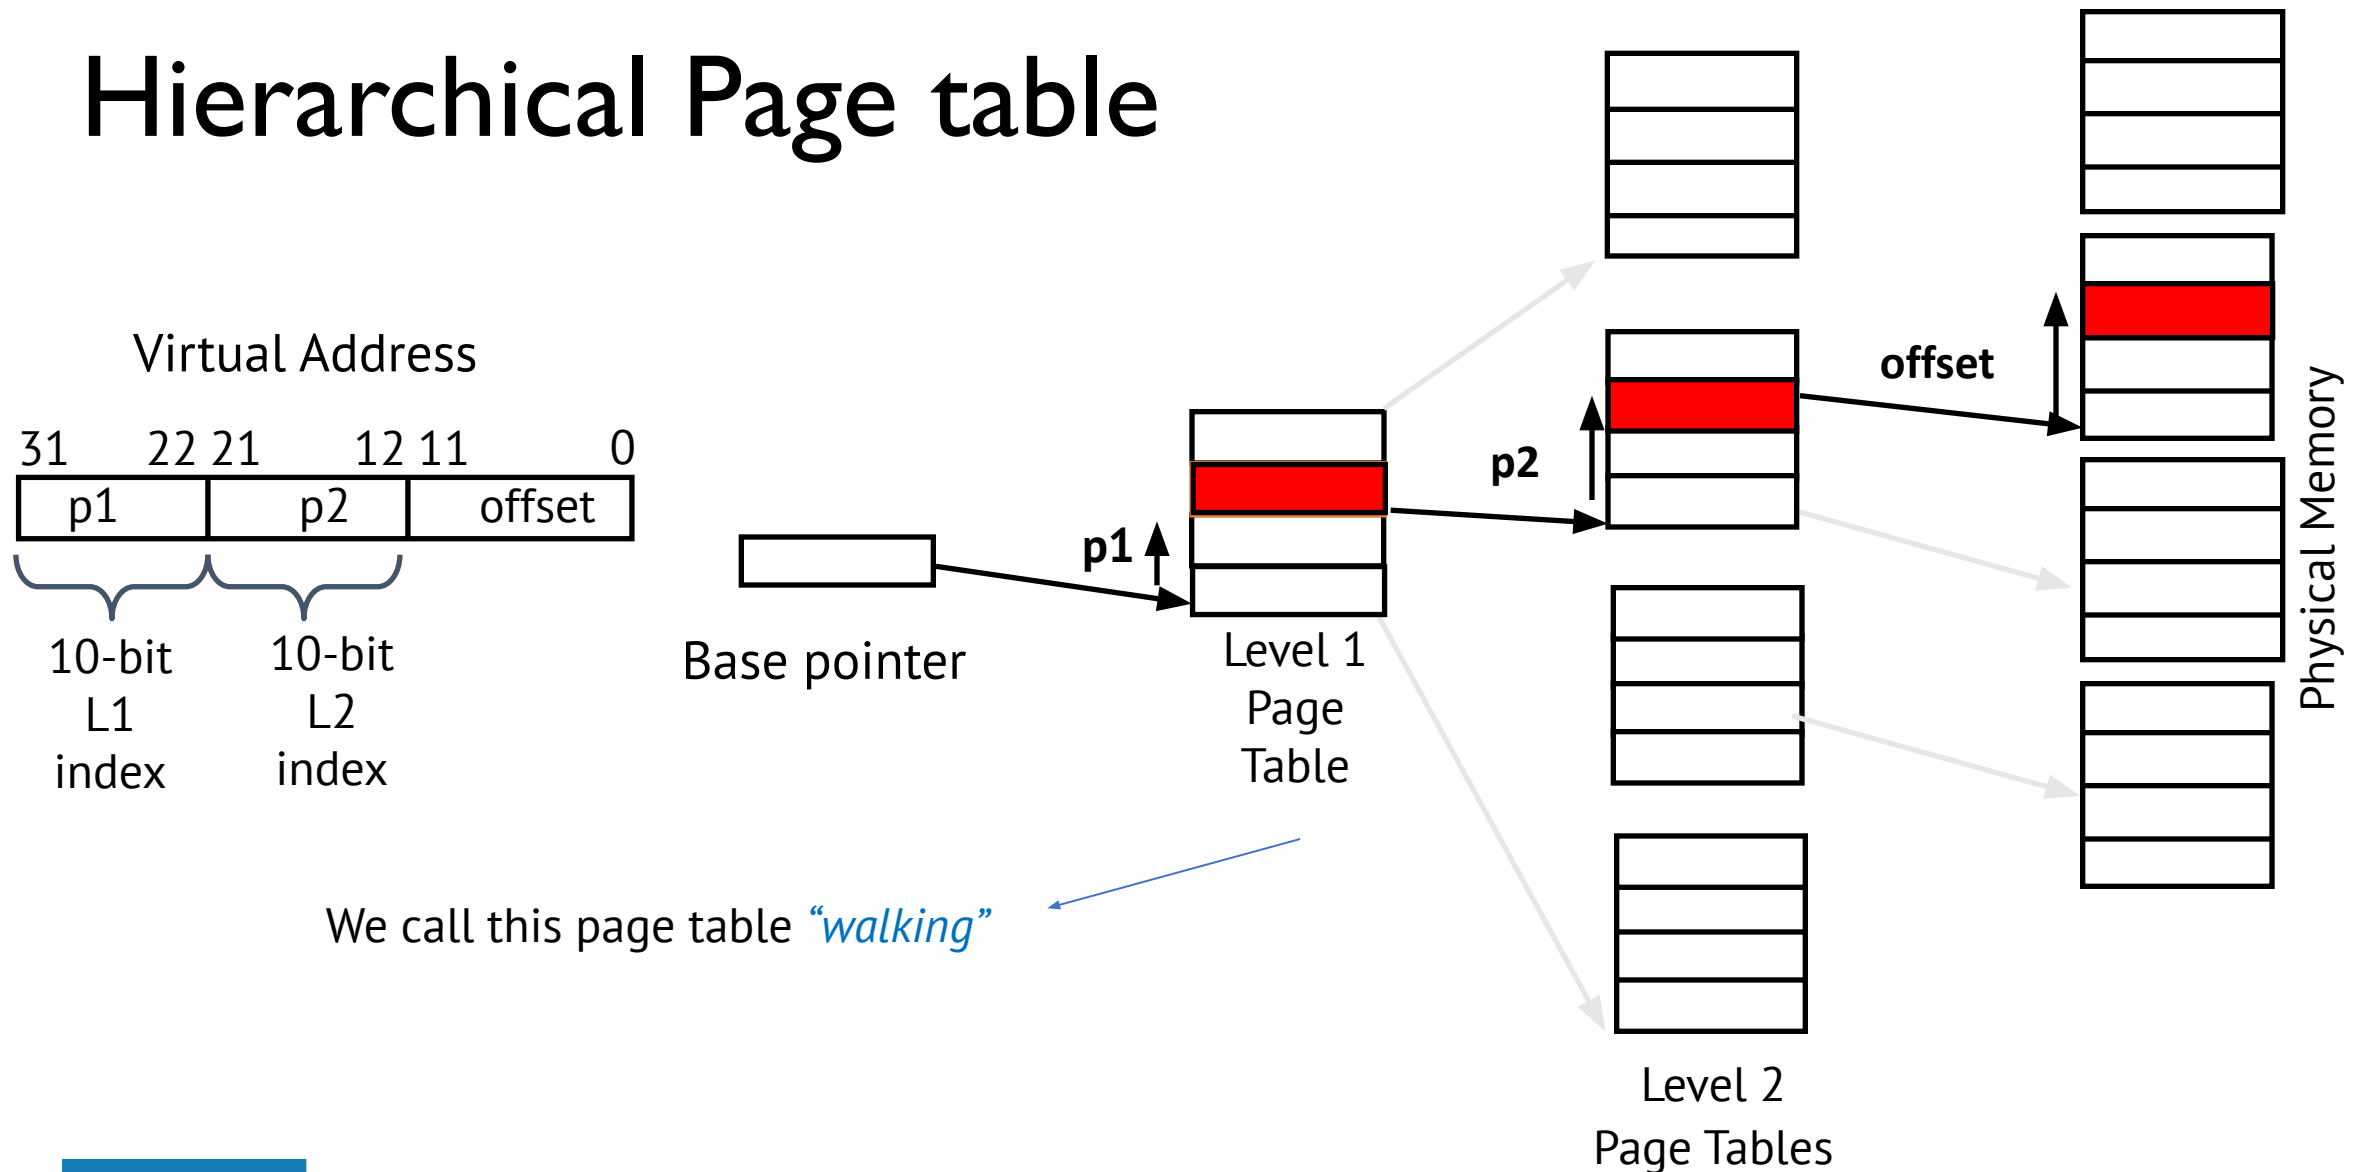
- Store the page table as a *page table*.



# Hierarchical Page table

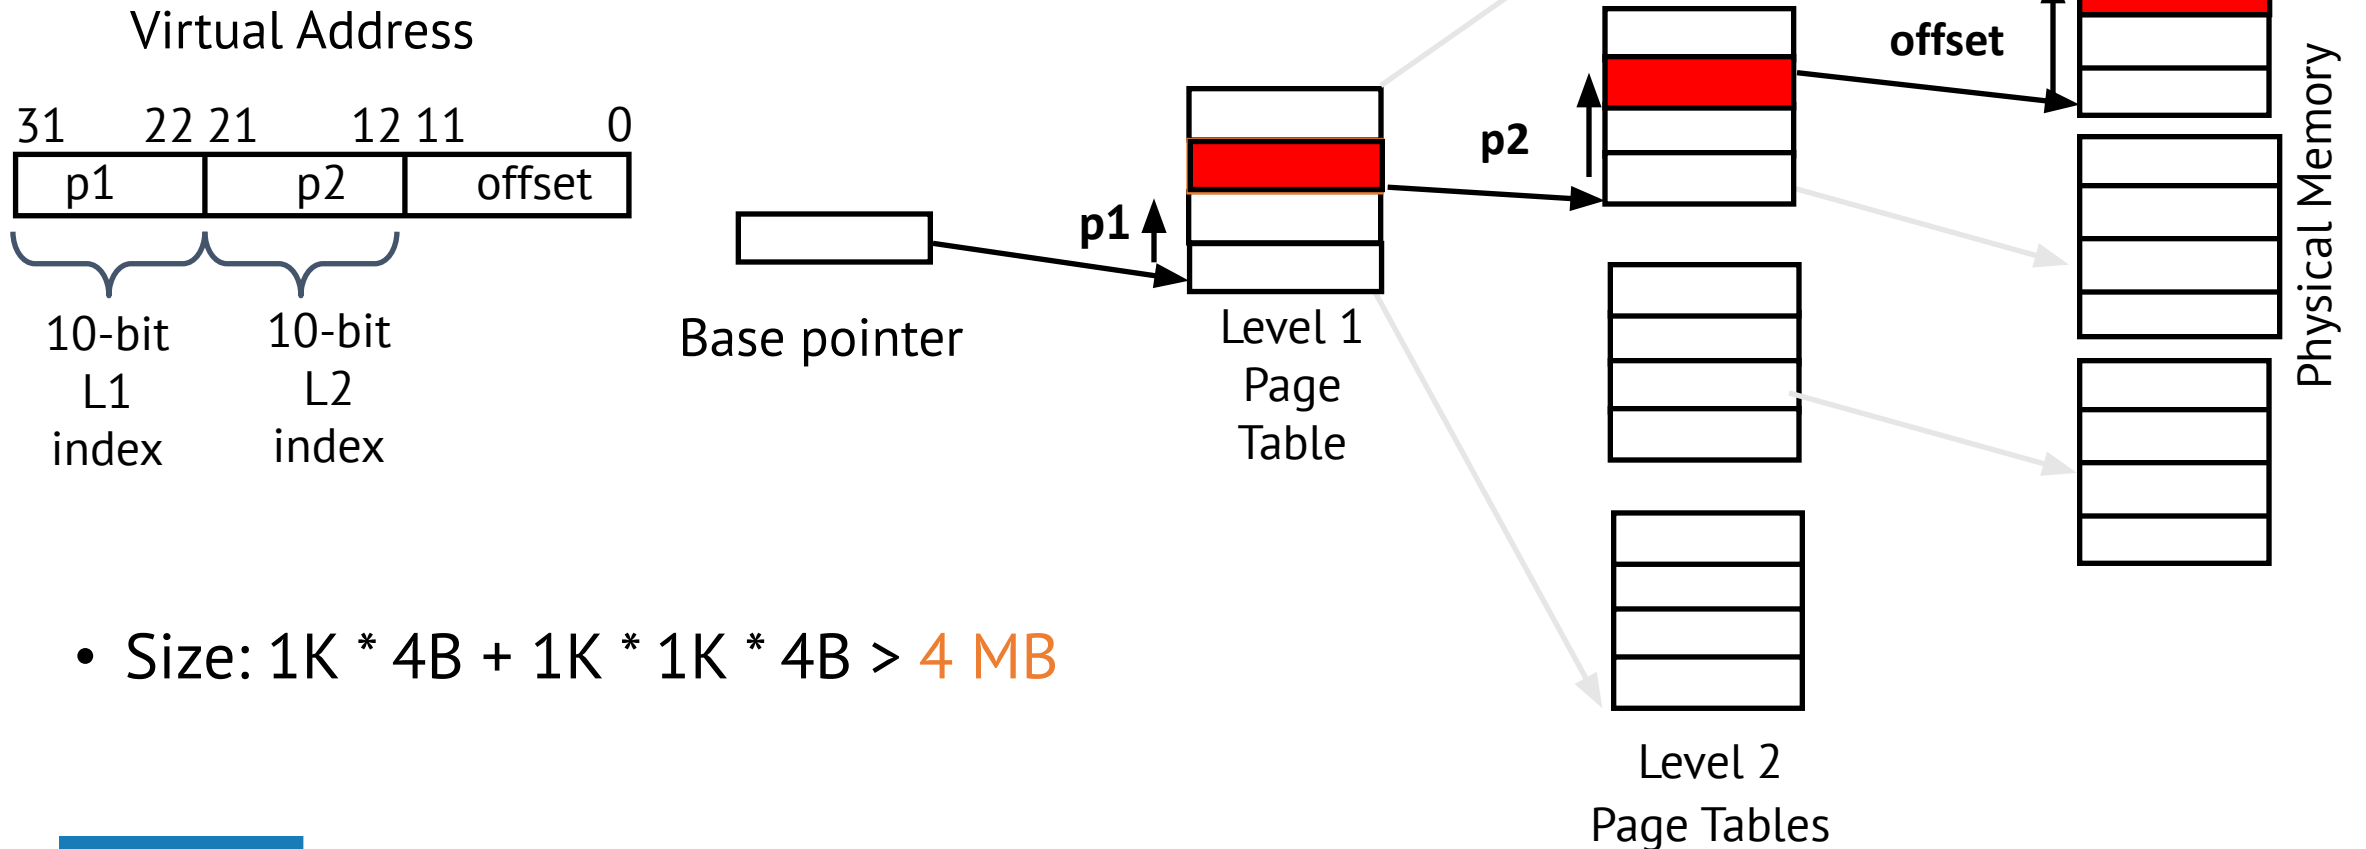


# Hierarchical Page table



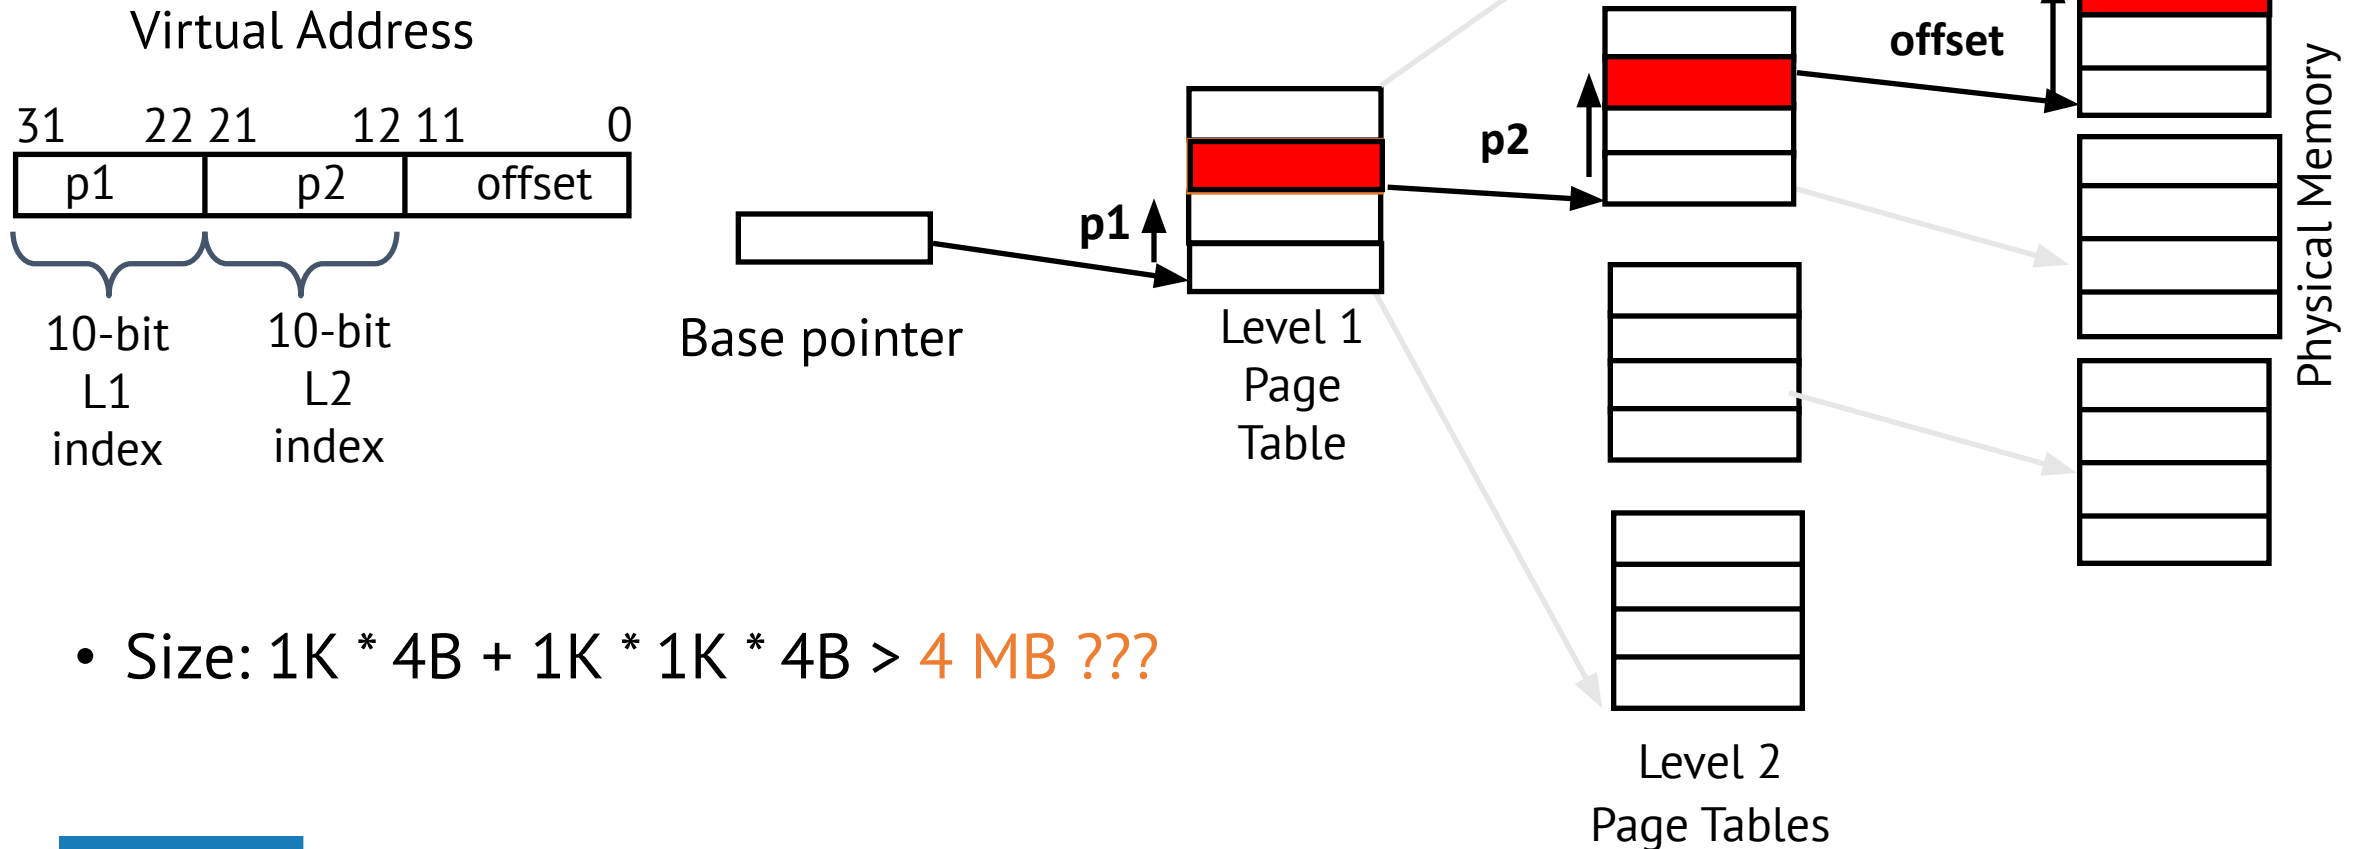


# Hierarchical Page table



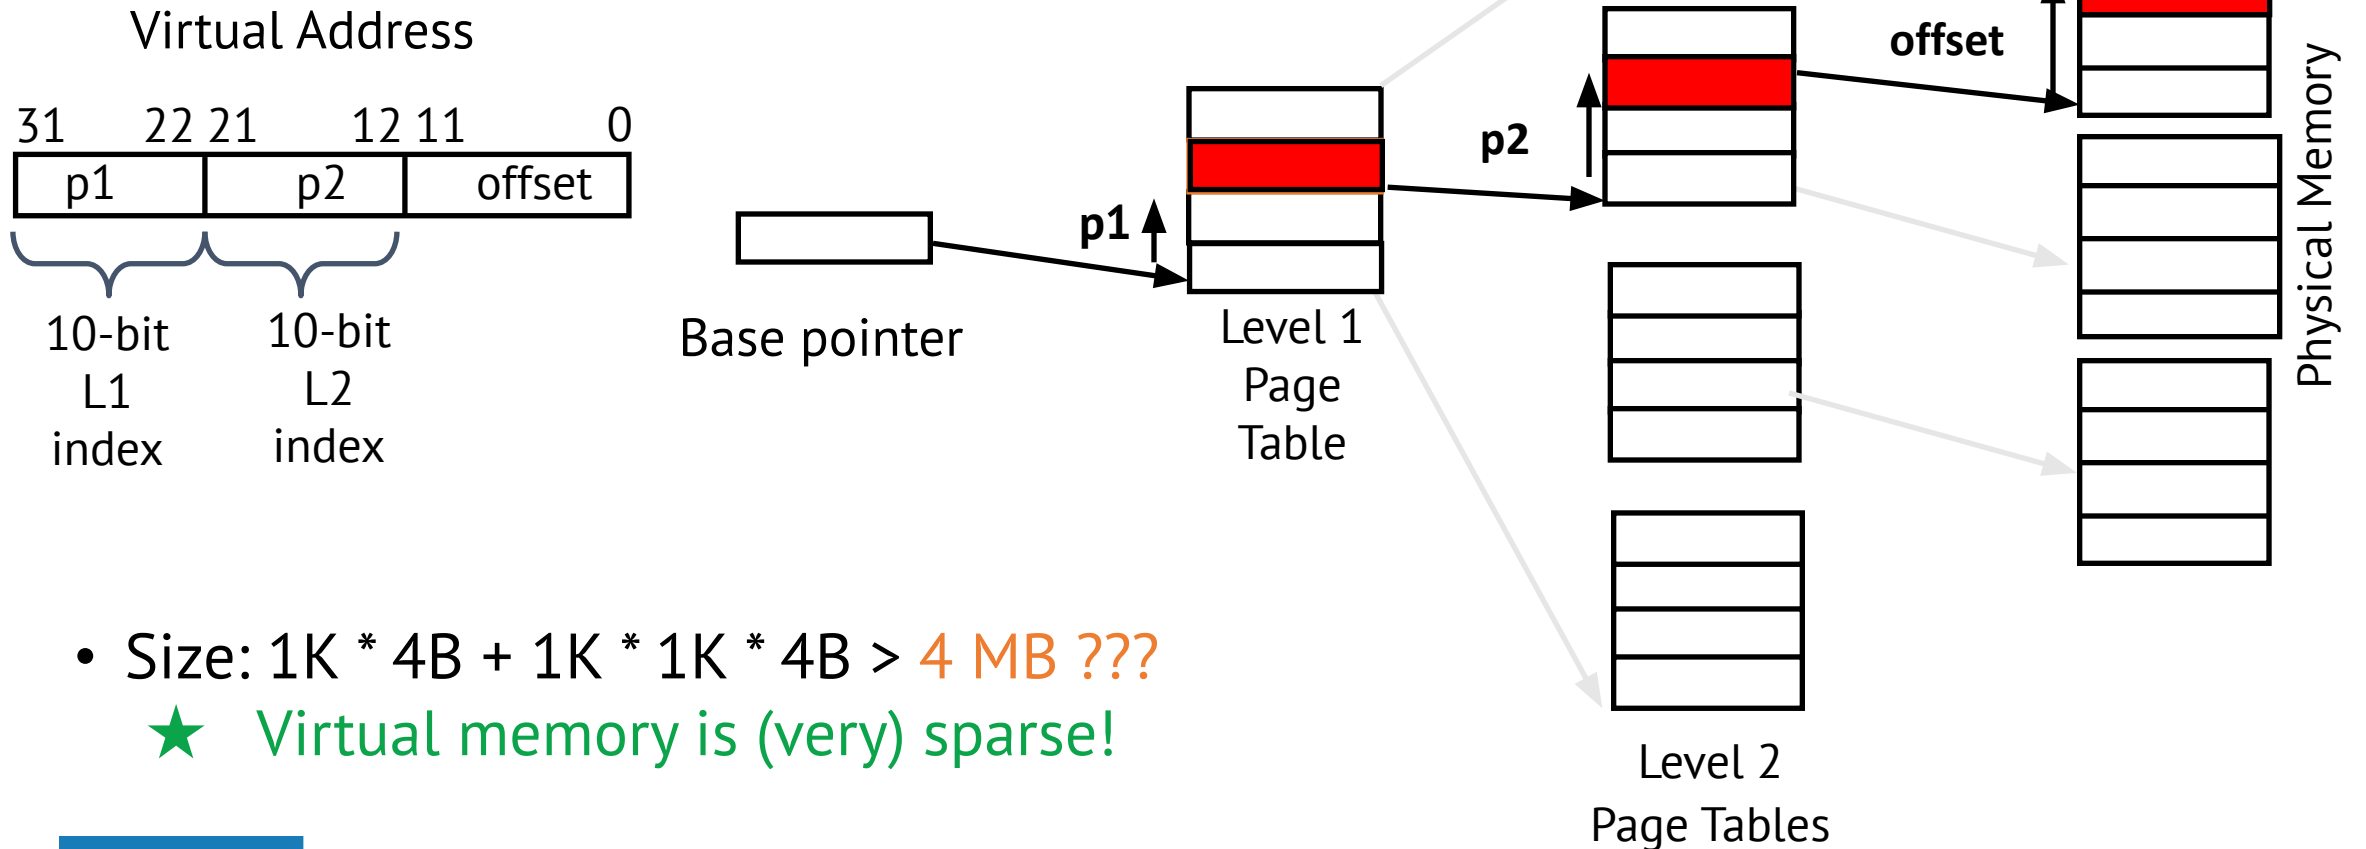
- Size:  $1K * 4B + 1K * 1K * 4B > 4 \text{ MB}$

# Hierarchical Page table



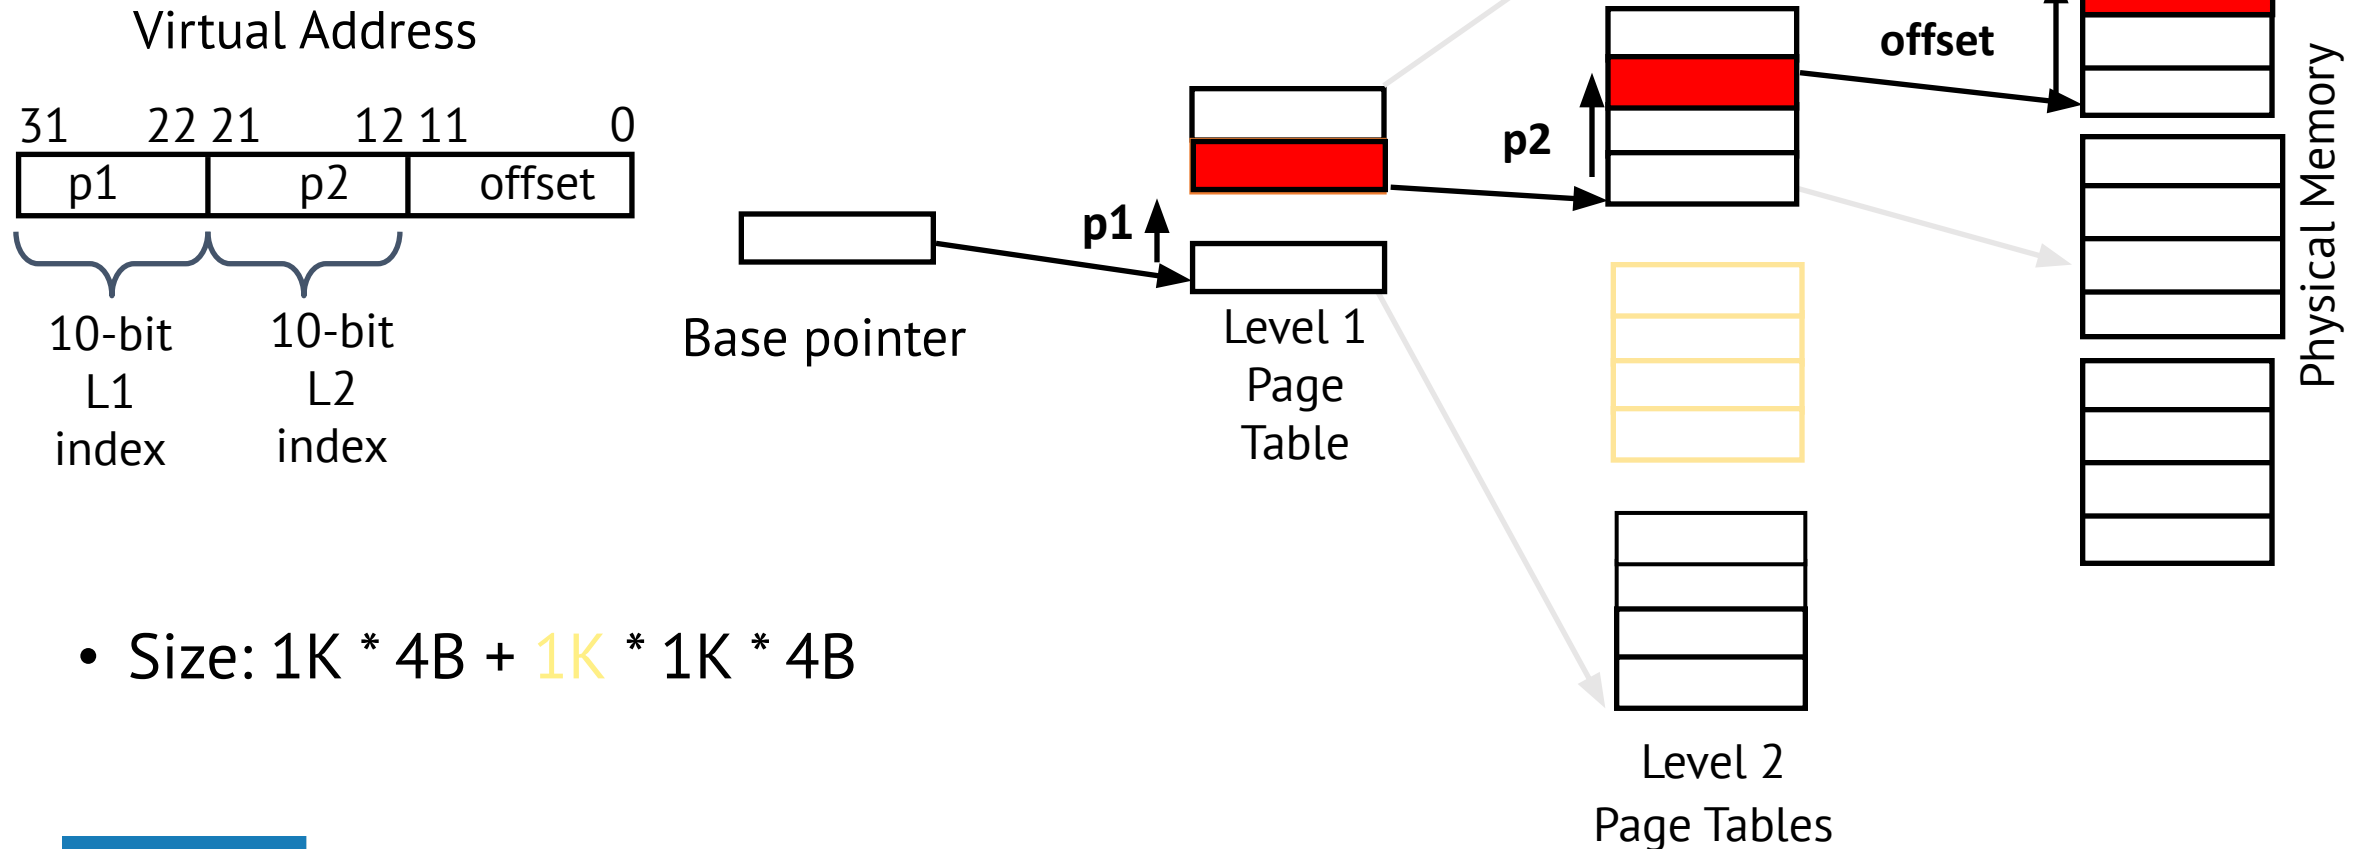
- Size:  $1K * 4B + 1K * 1K * 4B > 4 \text{ MB} ???$

# Hierarchical Page table



- Size:  $1K * 4B + 1K * 1K * 4B > 4 MB ???$   
★ Virtual memory is (very) sparse!

# Hierarchical Page table



- Size:  $1K * 4B + 1K * 1K * 4B$

# Hierarchical Page Tables

- We need page tables for address translation.
- Page tables are large and should be in the memory itself.
- Hierarchical Page tables can leverage the existing sparsity of virtual addresses and makes the page table storage compact!

# Memory Access

- Each memory access (e.g., L1 cache) needs **multiple** (three or more) memory accesses!

→ *What can we do?*

# Memory Access

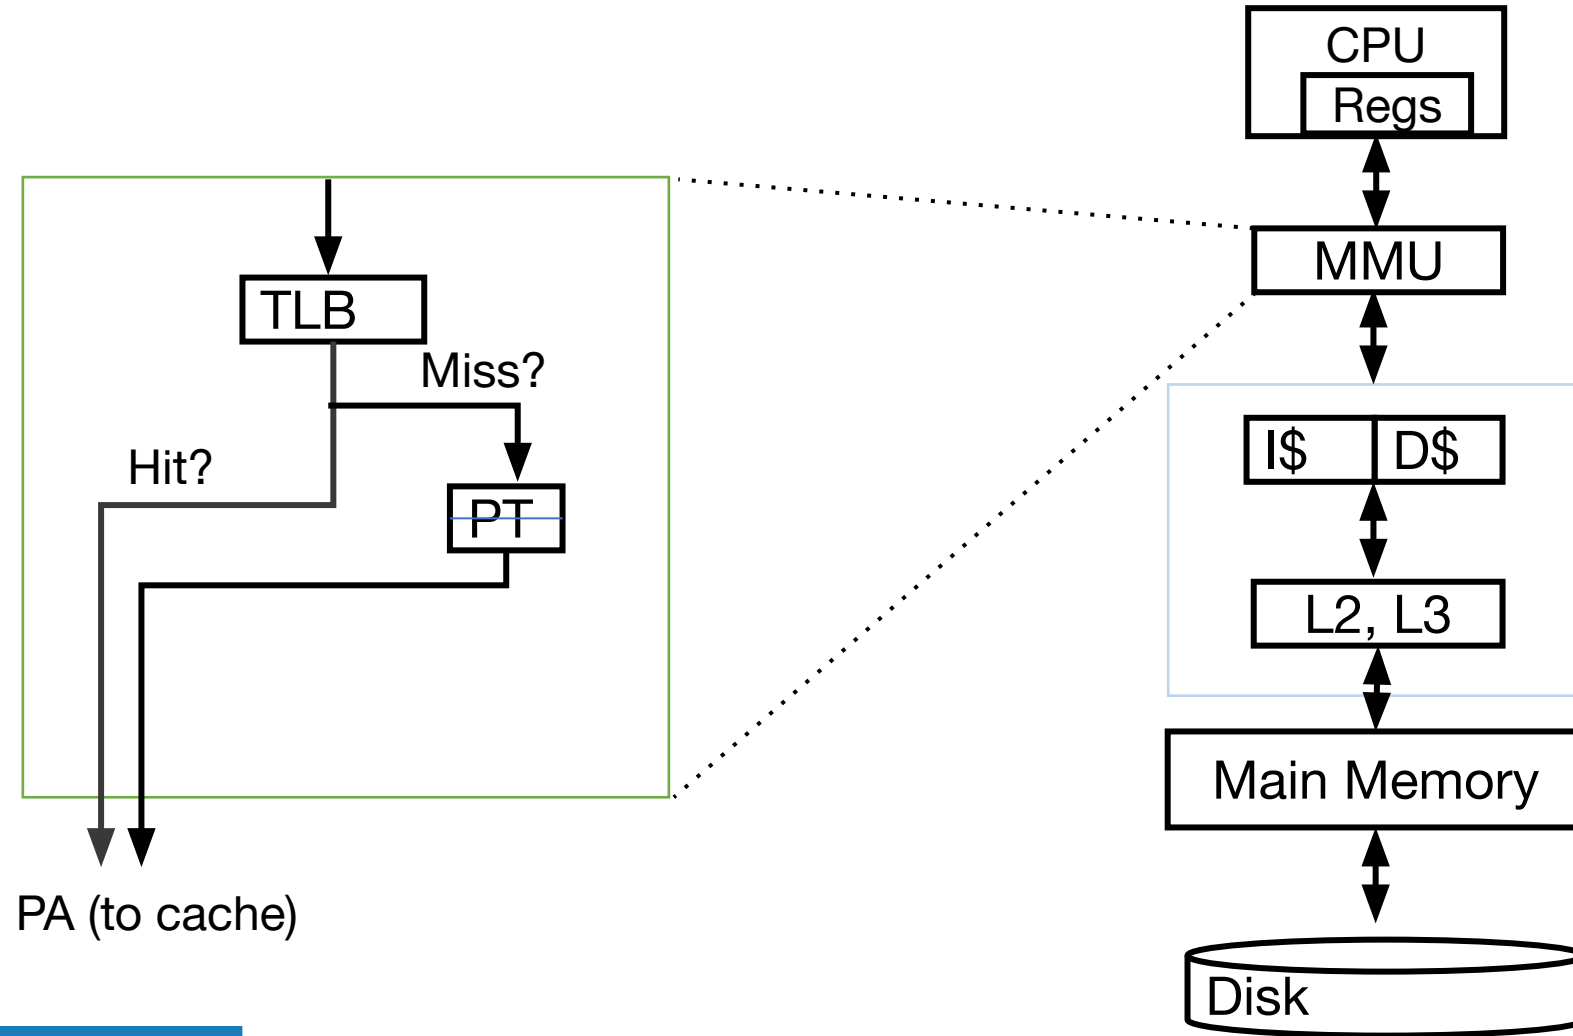
- Each memory access (e.g., L1 cache) needs **multiple** (three or more) memory accesses!

→ *What can we do?*

★ *Cache the most recent translations!*



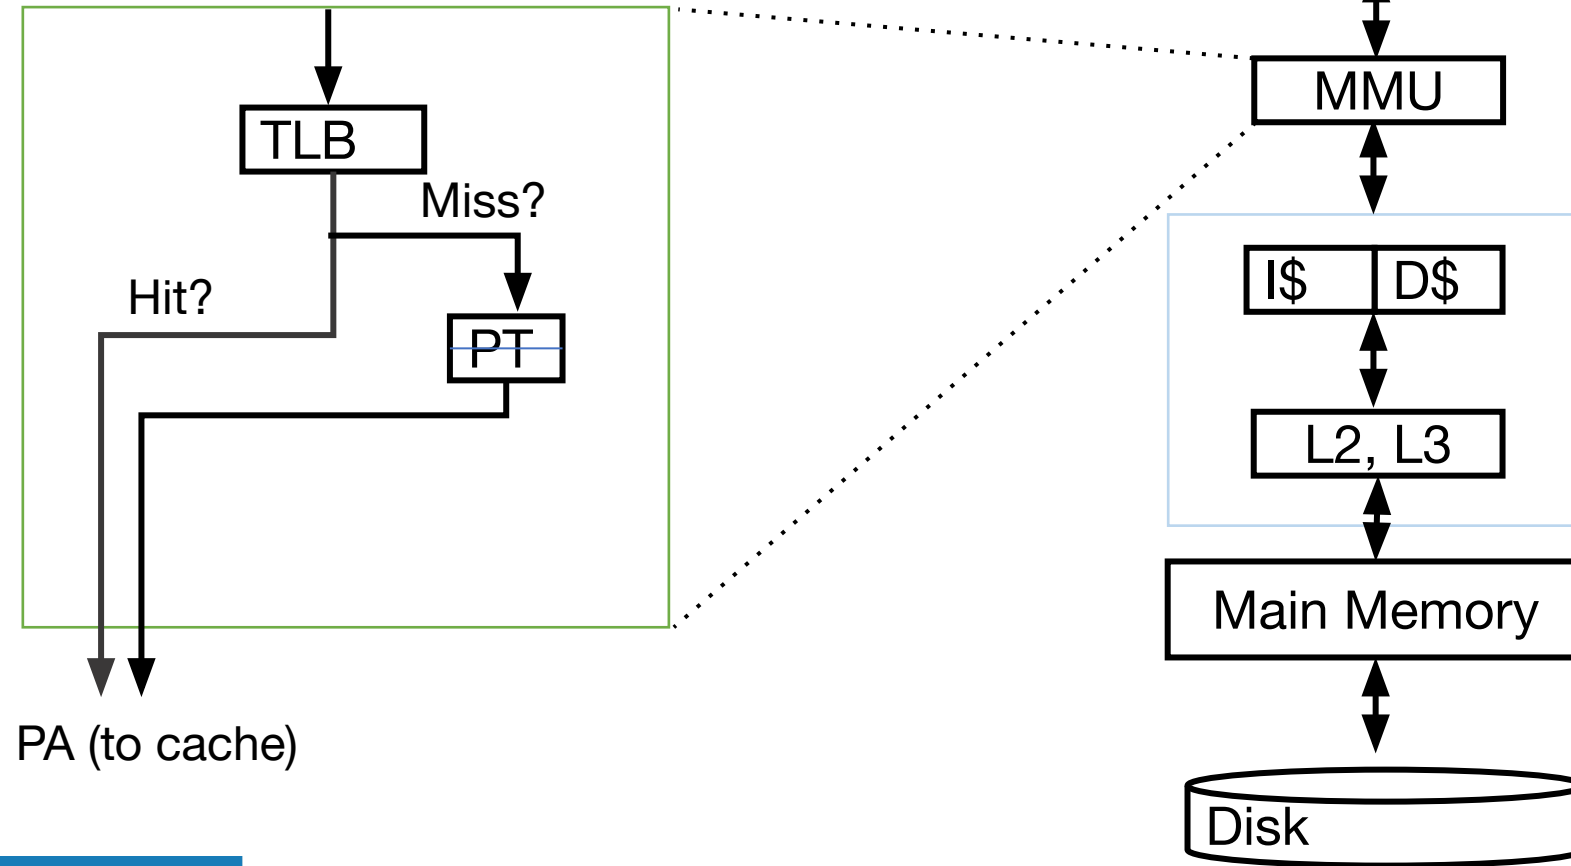
# Translation Lookaside Buffer (TLB)





# Hit time?

$$HT = t_{TLB} + t_{L1}$$

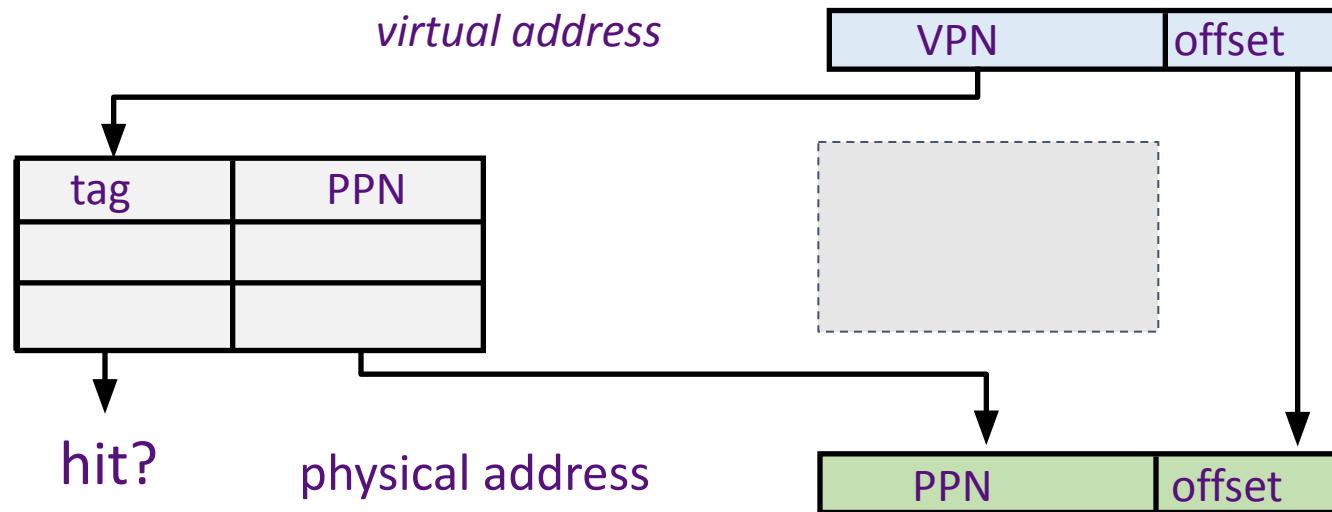


# Translation Lookaside Buffer (TLB)

- *Cache translations in TLB*

TLB hit  $\Rightarrow$  Single-Cycle Translation

TLB miss  $\Rightarrow$  Page-Table Walk to refill



# TLB Design

- Typically 32-128 entries, usually fully associative.
  - Each entry maps a large page, hence less spatial locality across pages.
  - Larger TLBs (256-512 entries) are 4 or 8 way set-associative.
  - Larger systems sometimes have multi-level (L1 and L2) TLBs.
  - Random or FIFO replacement policy.
  - **TLB Reach:**
    - Size of largest virtual address space that can be simultaneously mapped by TLB
    - **Example:** 64 TLB entries, 4KB pages, one page per entry  
 $TLB\ Reach = 256\ KB$

- TLB miss causes an exception and results in a page table walk.
- OS typically is responsible to handle TLB miss (software handling).
- Alternatively, memory management unit (MMU) can handle TLB miss.

# What about the Disk?



# How to utilize disk?

- Disk is our secondary storage unit with much bigger size, but much larger access time.

# How to utilize disk?

- Disk is our secondary storage unit with much bigger size, but much larger access time.

→ *How to efficiently use the disk?*

# Demand Paging

- ★ Use main memory and “swap” pages in the disk as automatically managed memory hierarchy levels.
  - Analogous to cache vs. main memory

# Demand Paging

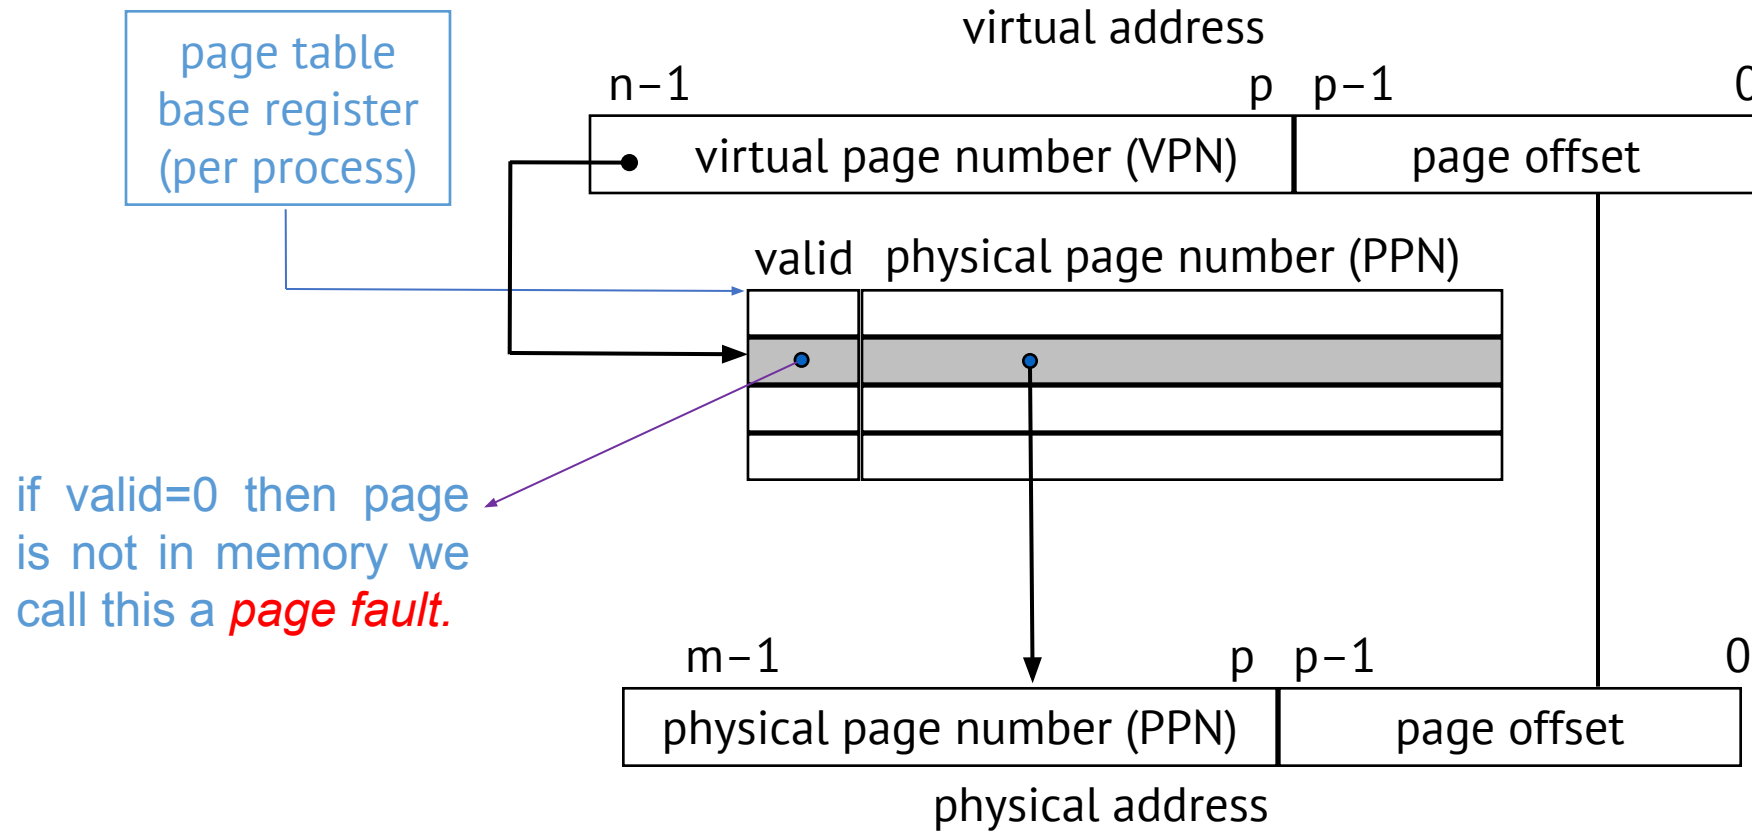
★ Use main memory and “swap” pages in the disk as automatically managed memory hierarchy levels.

→ Analogous to cache vs. main memory

- **M** (DRAM + Disk capacity) bytes of storage,
  - keep most frequently used **C** bytes in DRAM ( $C \ll M$ )
  - Keep the rest in disk.
  - If the page is not in DRAM, we call it a *page fault*.
  - Bring a page (from disk to main memory) when “demanded”.



# Page Table



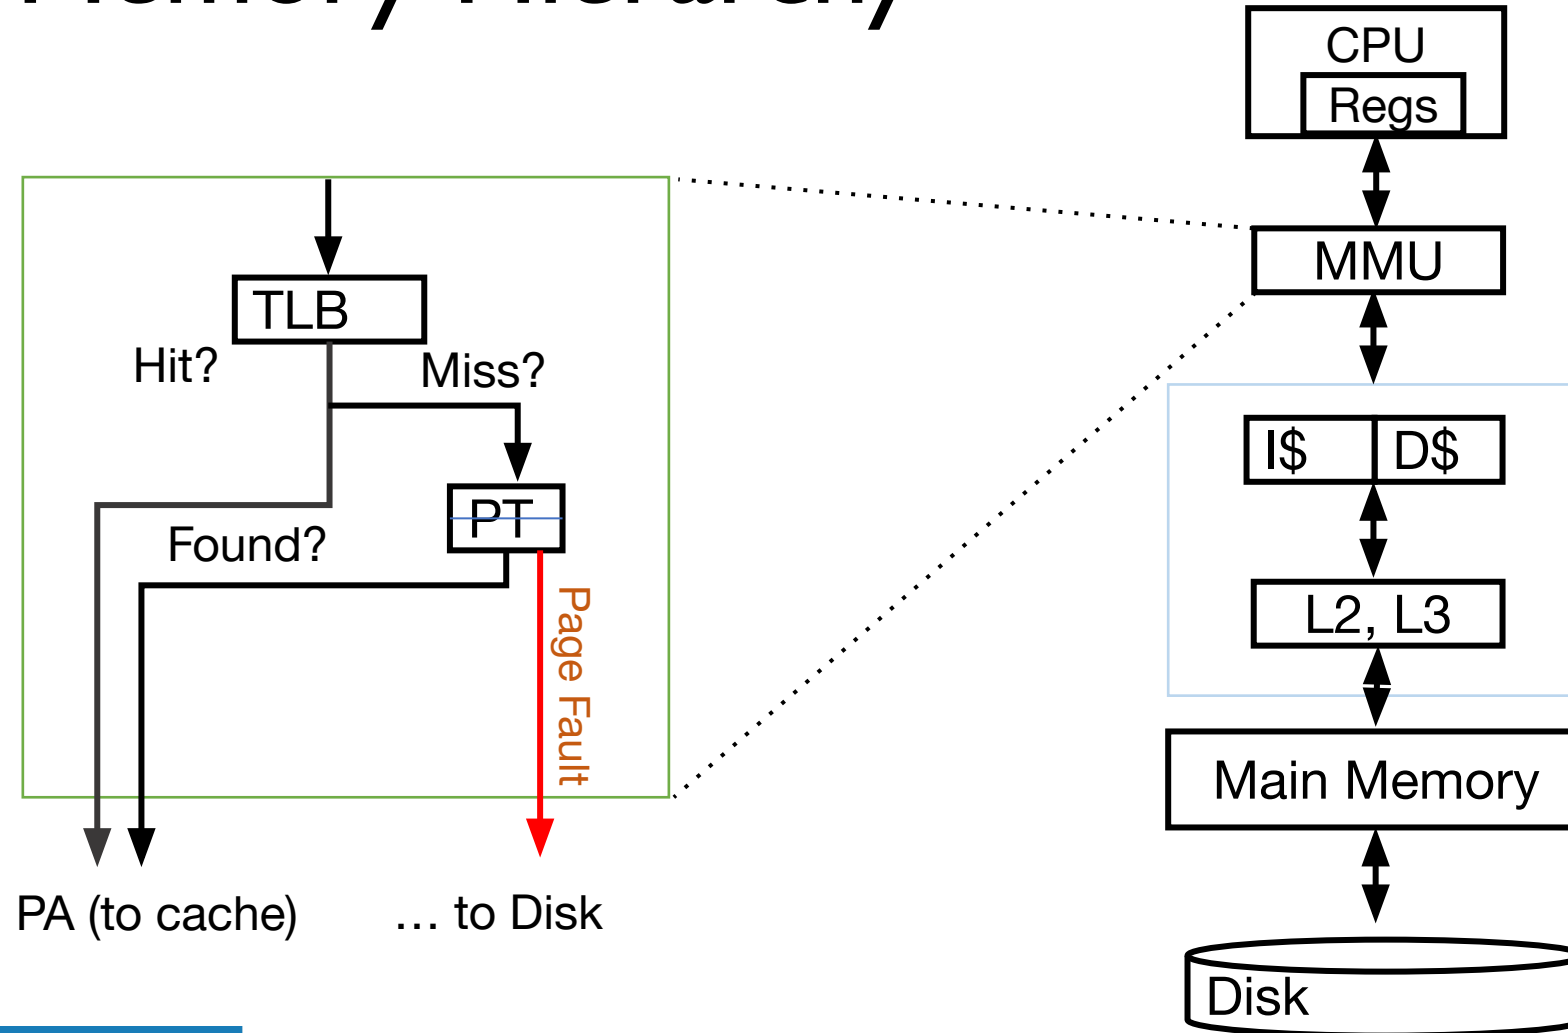
# Demand Paging Design

- *Same basic issues as before (in cache)*
  - When to bring a page into DRAM?
  - Which page to evict (we call it “swap”) from DRAM to disk to free-up DRAM for new pages?
  - Page Size?
  - ...

# Demand Paging Design

- *Same basic issues as before (in cache)*
- ★ *OS handles everything (easier, fast enough)*
  - Pseudo-LRU replacement policy

# Memory Hierarchy



# Recap

- Page tables
  - Why we need them and how to translate?
  - How to store them?
  - How to reduce access time?

# Recap

- Page tables
  - Why we need them and how to translate?
  - How to store them?
  - How to reduce access time?

→ Can we do better?

# How to decrease hit time?

- Can we access LI before TLB?

# How to decrease hit time?

- Can we access LI before TLB?
  - Address in LI should be stored with **VA** (i.e., no translation)!  
*Problem?*



# How to decrease hit time?

- Can we access LI before TLB?
  - Address in LI should be stored with **VA** (i.e., no translation)!

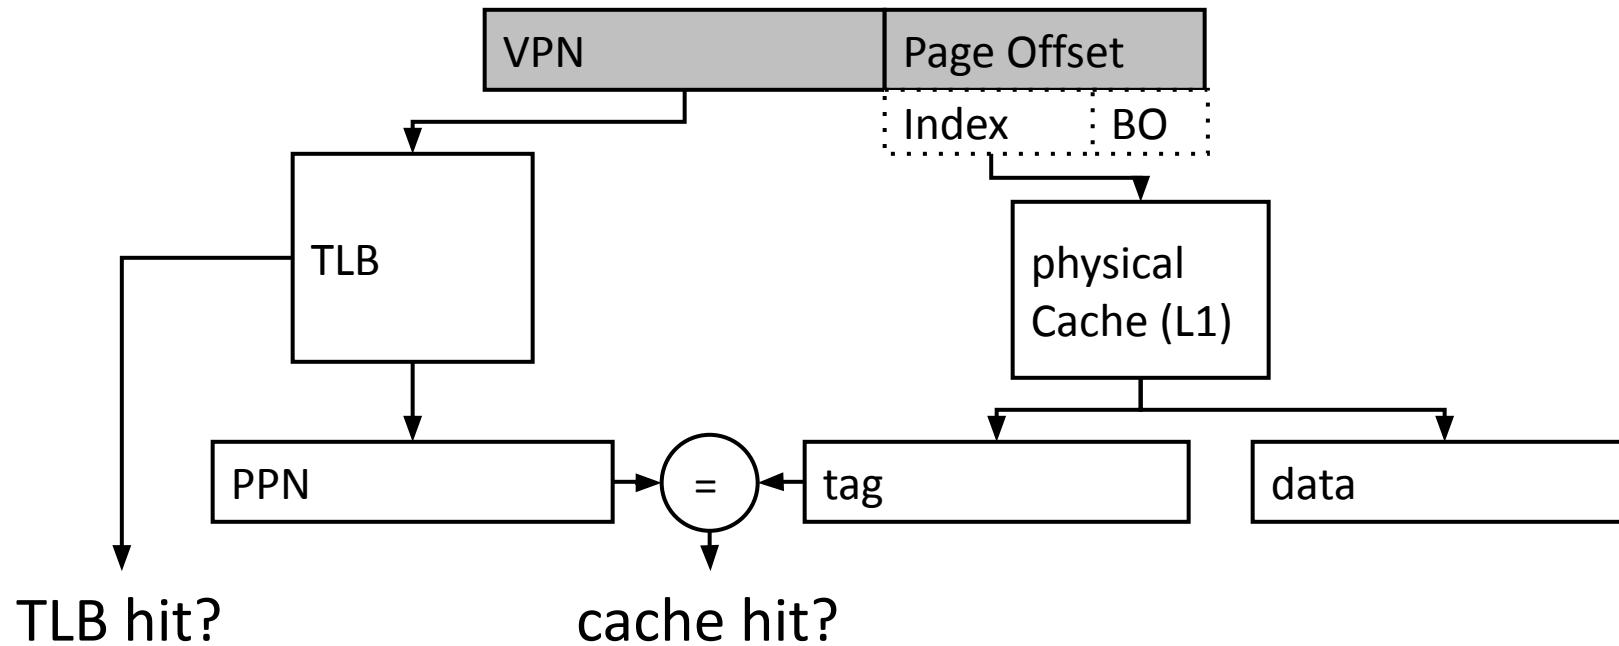
*Problem?*

- *Aliasing*

# How to decrease hit time?

- Can we access LI before TLB?
  - Address in LI should be stored with VA!  
- NO!
- Can we access TLB and LI *in parallel* ?

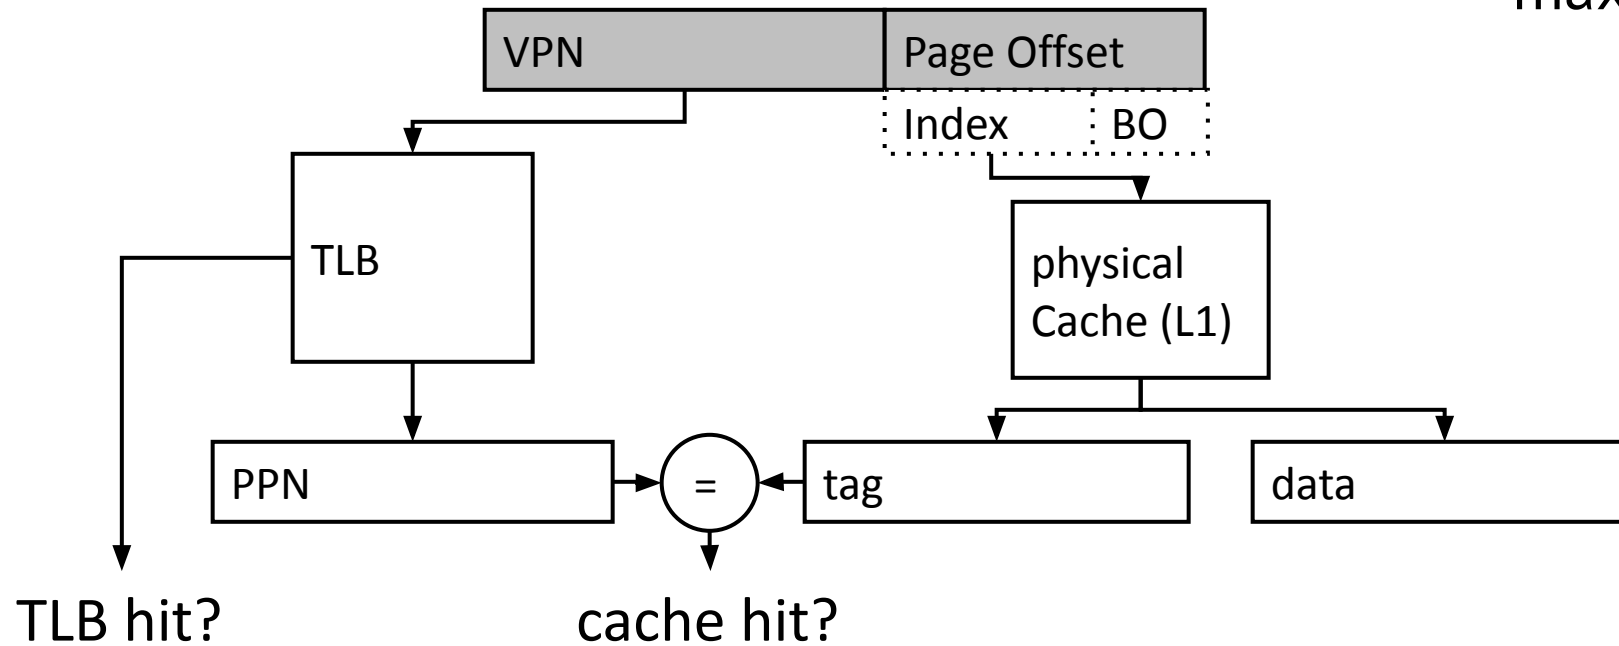
# *Virtually*-Indexed *Physically*-Tagged Cache



# *Virtually*-Indexed *Physically*-Tagged Cache

*Hit time?*

$$= \max(t_{\text{TLB}}, t_{\text{L1}}) = t_{\text{L1}}$$



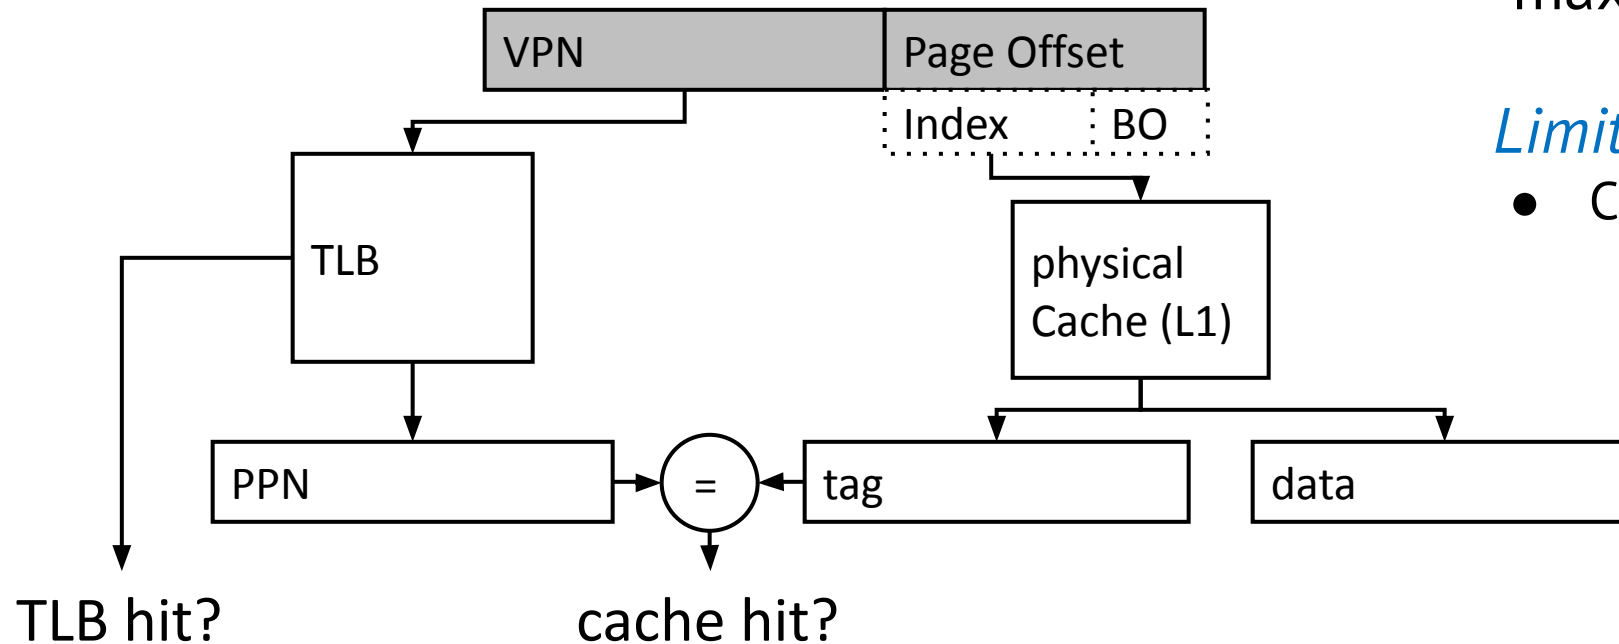
# *Virtually*-Indexed *Physically*-Tagged Cache

*Hit time?*

$$= \max(t_{\text{TLB}}, t_{\text{L1}}) = t_{\text{L1}}$$

*Limitation?*

- Cache size  $\leq$  Assoc. x Page Size



# With and Without MMU

- Most embedded processors and DSPs provide *physical addressing* only!
  - Can't afford area/speed/power budget for virtual memory support.
  - Often there is no secondary storage to swap to!

# End of Presentation



# Acknowledgement

- This course is partly inspired by the following courses created by my colleagues:
  - CSI52, Krste Asanovic (UCB)
  - I8-447, James C. Hoe (CMU)
  - CSE141, Steven Swanson (UCSD)
  - CIS 501, Joe Devietti (Upenn)
  - CS4290, Tom Conte (Georgia Tech)
  - 252-0028-00L, Onur Mutlu (ETH)