



Module 15 – Memory Consistency

(ECE M116C- CS M151B) Computer Architecture Systems

Nader Sehatbakhsh

Department of Electrical and Computer Engineering

University of California, Los Angeles

Next Problem: *Consistency!*



Consistency Problem

- if initially $M[Z] = M[Y] = 0$, $x3 = 100$, $x4 = 200$

Core1: lw x1, Z(x0)

Core2: lw x2, Y(x0)

...

...

Core1: sw x3, Y(x0)

Core2: sw x4, Z(x0)

→ What are the potential values for x1 and x2?

Wait, but why?

- RAW?
- Correctness?
- LSQ?
 - If the addresses are not the same, why wait?
 - Compiler can reorder, too!

Consistency Problems

- SW and LW (to different addresses) within one core can be reordered and that will mess up assumptions in OTHER cores.
- The timing between cores might be off! (even without reordering).

Memory Consistency vs. Coherency

- Consistency is about ordering of parallel accesses between *different* addresses.
- Coherency is about ordering of parallel accesses to the *same* address.

So, how to fix this?

The Timing Problem

- Timing between two cores is not the hardware or reordering problem, so it's programmer's job to fix that using synchronization strategies.
- We will discuss this later!

The Reordering Problem

- That is the optimization done by hardware and/or compiler and programmer cannot do anything about this!
- We need to fix this!

Memory Model

- We first need to know what are the possibilities.

→ **Memory Model:** Interactions between threads and the shared memory. In other words, *what are the possible values for a load?*

Sequential Consistency

SC: *what programmers expected.*

1. Processors see their own loads and stores in program order.
2. Processors see others' loads and stores in program order.
3. All processors see the same global load/store ordering.

Sequential Consistency

- Simplest scenario!
- SC makes multicore system indistinguishable from multi-programmed in-order uniprocessor.

Sequential Consistency

- Simplest scenario!
- SC makes multicore system indistinguishable from multi-programmed in-order uniprocessor.
- **Too slow!**
 - Too many restrictions (what are they?)



Memory Operation Ordering

- Four types of memory operation orderings
 - $W \rightarrow R$: write must complete before subsequent read.
 - $R \rightarrow R$: read must complete before subsequent read.
 - $R \rightarrow W$: read must complete before subsequent write.
 - $W \rightarrow W$: write must complete before subsequent write.

- This is for *different* addresses, so no RAW hazards!

What can we do?

- We really need to be able to re-order memory accesses.
 - Most lines are not even shared!
 - Store does not need to even wait
 - We have the ability to do so!

What can we do?

★ We need to compromise, and use a “weaker” or “relaxed” model.

→ *Not all loads and stores need to be in-order.*

Memory Operation Ordering


- Four types of memory operation orderings
 - $W \rightarrow R$: write must complete before subsequent read.
 - $R \rightarrow R$: read must complete before subsequent read.
 - $R \rightarrow W$: read must complete before subsequent write.
 - $W \rightarrow W$: write must complete before subsequent write.
- *Sequential consistency maintains all four orderings.*
- *Relaxed memory consistency models allow **certain** orderings to be violated.*

Total Store Order Model



Total Store Order (TSO) Model

- Processor P can read B before its write to A is seen by all processors.
 - Each processor can move its own reads in front of its own writes.
 - *Remember LSQs?*



Total Store Order (TSO) Model

- Processor P can read B before its write to A is seen by all processors.
 - Each processor can move its own reads in front of its own writes.
- NOTE: Read by other processors *cannot* return new value of A until the write to A is observed by all processors.

How TSO helps?

`a = 0, flag = 0`

thread 1

```
store 1 → a  
store 1 → flag
```

thread 2

```
loop: if (flag == 0) goto loop  
      load a
```

→ *Can load read 0?*

TSO vs. SC

- if initially $M[Z] = M[Y] = 0$

Core1: lw x1, Z

...

Core1: sw Y, 100

Core2: lw x2, Y

...

Core2: sw Z, 200

→ What are the potential values for x1 and x2?

TSO Summary

- Improving the per-core performance by allowing the reordering within the core!
- Still not observable by other cores (timing problem)

Weak Consistency

- Different assumptions
 - What and when a memory access is visible to the other cores.
- Performance vs. Complexity
- Different architectures use different models.
- More about it soon!

RISC-V Memory Model

- RISC-V OoO (BOOM) follows the RVWMO memory consistency model (RISC-V weak memory ordering).

RISC-V Memory Model

- Loads are optimistically fired to memory on arrival to the LSU (for performance).
- Simultaneously, the load instruction compares its address with all of the store addresses that it depends on.
 - If there is a match, the memory request is killed.
 - If the corresponding store data is present, then the store data is forwarded to the load and the load marks itself as having succeeded.
 - If the store data is not present, then the load goes to sleep.
 - Loads that have been put to sleep are retried at a later time.

RISC-V Memory Model

- RISC-V OoO (BOOM) follows the RVWMO memory consistency model (RISC-V weak memory ordering).
- BOOM currently exhibits the following behavior:
 - Write → Read constraint is relaxed (newer loads may execute before older stores).
 - Read → Read constraint is maintained (loads to the same address appear in order).
 - A thread can read its own writes early.

How to write a correct program?

- The memory model is given.

Simple Solution

- Let it be!
 - Race-free programming.



Not very simple solution

- Have the hardware to track potential violations!
 - Each core reorders but if another core needs that data, it should recover.
 - Very costly!


More realistic solution

- Use synchronization!
 - The programmer (and/or hardware) use some primitives to synchronize different programs.
 - This solves both timing and ordering problems.

Synchronization

- Two major solutions:

I - Producer-consumer model




Synchronization

- Two major solutions:

1- Producer-consumer model

2- Mutual exclusiveness




Solution

- Use *fences*!

- FENCE is an instruction that enforces all memory operations *before* it to complete before fence is executed.
- All memory operations *after* FENCE must wait for FENCE to be finished.
- This happens *within* each core!

FENCE in RISC-V

- FENCE (rwio)
 - fence r,r



What about between cores?

Semaphore

- Create some barriers!
- This will synchronize all the processors.
- Typically implemented in software.

Synchronization Using Semaphores

Core1: lw x1, Z

...

Core1: sw Y, 100

Core2: lw x3, Y

...

Core2: sw Z, 200

How to implement barrier/semaphore in hardware?

- Hardware-or with one-hot encoding → set a flag to one when a core reaches to the barrier. Wait until all flags are one!

What about mutual exclusiveness?

Mutual Exclusiveness

- Making sure only one process has access to the shared region at any given time.

Mutual Exclusiveness

- Making sure only one process has access to the shared region at any given time.

→ How?

How to fix interleaving?

- Use Mutex/Locks!

Acquire (lock)

Core1: lw x1, Z

...

Core1: sw Y, 100

Release (lock)

Acquire (lock)

Core2: lw x3, Y

...

Core2: sw Z, 200

Release Consistency

- It is guaranteed that all reads and writes are executed and visible to all cores when the lock is released.
- Other than this guarantee, any ordering outside the acquire/release blocks is possible.

How to implement acquire/release?

```
acquire(lock) {  
    while (lock != 0)  
    {} // busy wait  
    // lock is acquired at this point  
    lock = 1; // making sure no one else can get it  
}
```

Assembly Version

loop:

```
lw ra, 0(x1)
```

```
bnez ra, loop
```

```
addi ra, x0, 1
```

```
sw ra 0(x1)
```

- Assume that Mem[x1] is where lock is stored.

How to implement acquire/release?

```
release(lock) {  
    lock = 0;  
}
```

- **Assembly Version:**

```
st r0 -> 0(&lock)
```

Problem?

(Core 1)

loop:

lw ra, 0(x1)

bnez ra, loop

addi ra, x0, 1

sw ra 0(x1)

(Core 2)

loop:

lw ra, 0(x1)

bnez ra, loop

addi ra, x0, 1

sw ra 0(x1)

Lock and Release

- Makes sense but we need to implement it such that locks are exclusives (i.e., only one lock is active at any given time).

Atomic Instructions

- A hardware primitive that guarantees instructions are executed in-order and without interruption.

How to implement Mutex?

- Hardware-enforce instruction (called test and set).

```
TS(int x) {  
    oldval = SWAP(x,1); // atomic swap  
    return oldval;  
}
```

- Other variants: *compare-and-swap, fetch-and-add*

How SWAP is implemented?

- Load-reserved, store-conditional (in RISC-V)

Label:

```
load-link x2, 0(x1)
```

```
addi x2, x0, 1
```

```
store-conditional x2, 0(x0)
```

```
branch-not-zero label // check for failure
```

LR and SC

- On load-link, processor remembers address.
- Then checks for writes by other processors.
 - If write is detected, store-conditional will fail.
- For MESI:
 - Load-Reserved ensures line in cache in Exclusive/Modified state.
 - Store-Conditional succeeds if line still in Exclusive/Modified state

Recap

- Locks

- Test-and-set
- Atomic
- LR and SC

→ This can be combined with a very weak consistency model (called release consistency).

Optimizations?

- Too many SWs when waiting/spinning.
→ Use test-test-and-set

Test-test-and-set

```
TTS(int x) {  
    oldval = x; // the read of x  
    if (oldval == 0) oldval = SWAP(x, 1);  
    return oldval;  
}
```

→ *while (TTS(lock) == 1);*

Optimizations?

- Too many SWs when waiting/spinning.
→ Use test-test-and-set
- Unfairness?
→ Use queues
- Multiple/Parallel Locks

Bottom-Line

- Range of possibilities!
- The burden is (mostly) on the developer!

CA 3

- Designing a MOESIF cache coherence protocol.

End of Presentation

Acknowledgement

- This course is partly inspired by the following courses created by my colleagues:
 - CSI52, Krste Asanovic (UCB)
 - I8-447, James C. Hoe (CMU)
 - CSE141, Steven Swanson (UCSD)
 - CIS 501, Joe Devietti (Upenn)
 - CS4290, Tom Conte (Georgia Tech)
 - 252-0028-00L, Onur Mutlu (ETH)